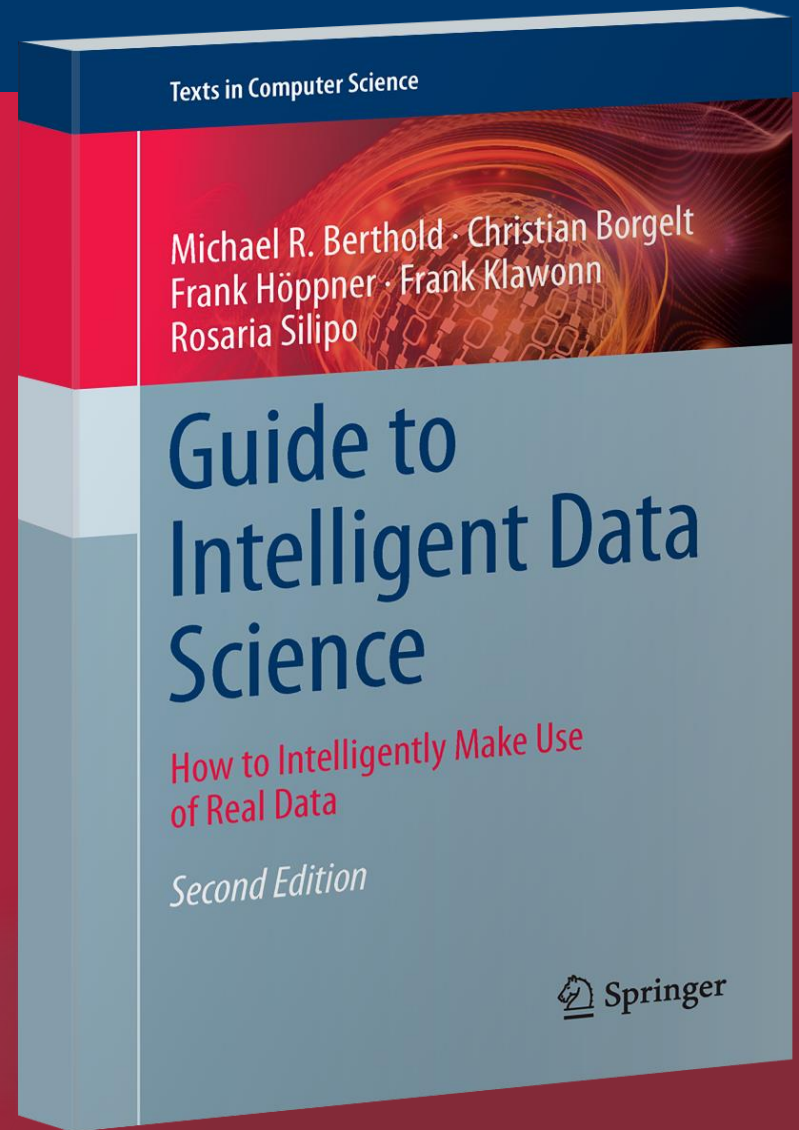


# Deep Learning



“What people call *AI* is no more than finding answers to questions we know to ask. Real *AI* is answering questions we haven't dreamed of yet”  
-*Tom Golway*

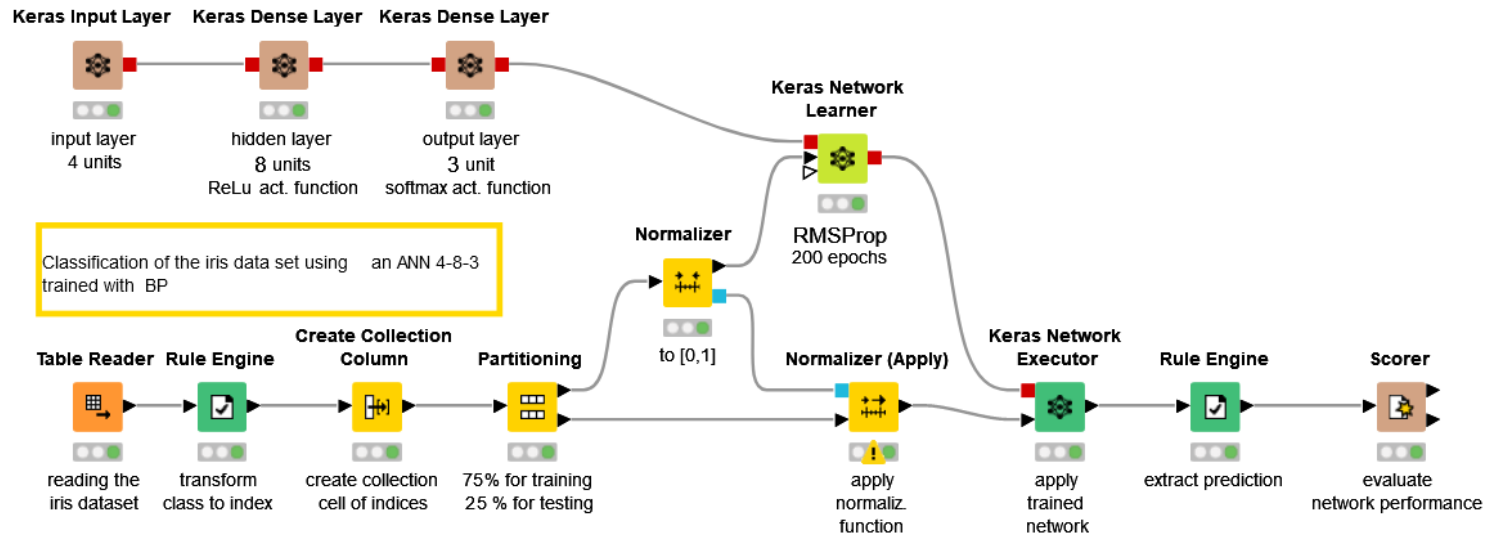
How deep can we dig in AI?

*\*This lesson refers to chapter 9 of the GIDS book*

## Content of this Lesson

- Recurrent Neural Networks (RNNs)
- Long Short Term Memories (LSTMs)
- Convolutional Neural networks (CNNs)
- Generative Adversarial Networks (GANs)

- Datasets used : iris dataset
- Example Workflows:
  - „Classifying the iris data set with ANN“ <https://kni.me/w/ei3eX9Sj5-RFEUat>
  - Keras layers
  - Multi-layer perceptron
  - Back propagation



- Deep Learning is the recent evolution of Neural Networks
- It covers:
  - Feedforward networks with many hidden layers (deep 😊)
  - New paradigms, like LSTMs in Recurrent Neural Networks, suitable for time series analysis
  - New topological layers, like convolutional and pooling layers, mainly for image processing
  - New architectures as in Generative Adversarial Networks (GANs)
  - ...
- Improvements are mainly due to:
  - Increased computational power for faster calculations, like GPUs
  - Parallel Computation

# Recurrent Neural Networks (RNNs)

## What are Recurrent Neural Networks?

- **Recurrent Neural Networks (RNNs)** are a family of neural networks suitable for processing of sequential data
- RNNs include auto and backward connections
- RNNs are used for all sorts of tasks:
  - Language modeling / Text generation
  - Text classification
  - Neural machine translation
  - Image captioning
  - Speech to text
  - Numerical time series data, e.g. sensor data
  - Time series analysis
  - ...

## Why do we need RNNs for Sequential Data?

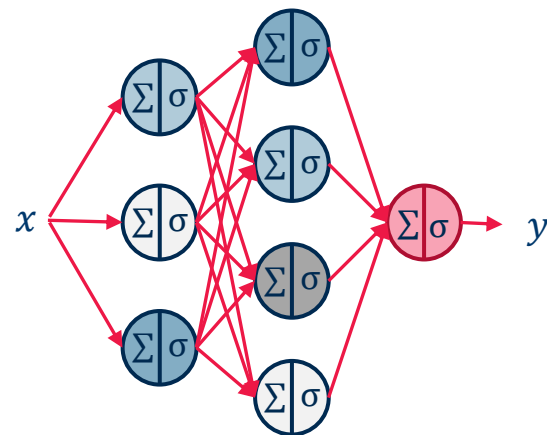
- **Goal:** Translation from German to English

*“Ich mag Schokolade”*  
 $\Rightarrow$  *“I like chocolate”*

- Option One: Use feed forward network to translate word by word

- But what happens with this question?

*“Mag ich Schokolade?”*  
 $\Rightarrow$  *“Do I like chocolate?”*

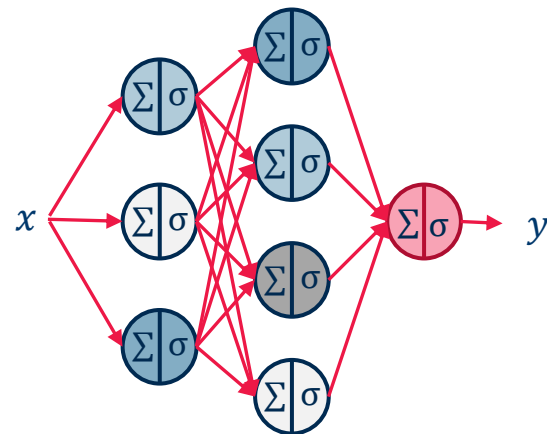


Input x	Output y
Ich	I
mag	like
Schokolade	chocolate



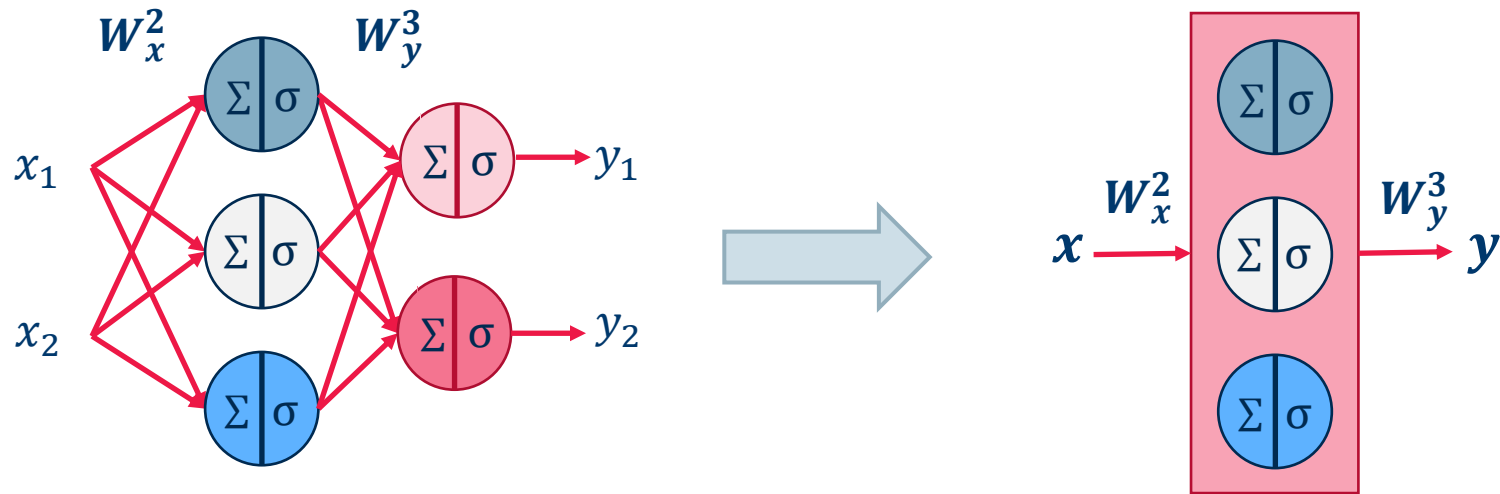
## Why do we need RNNs for Sequential Data?

- Problems with FFNN:
  - Each time step is completely independent
  - For translations we need context
  - More general: we need a network that remembers inputs from the past
- **Solution:** Recurrent Neural Networks

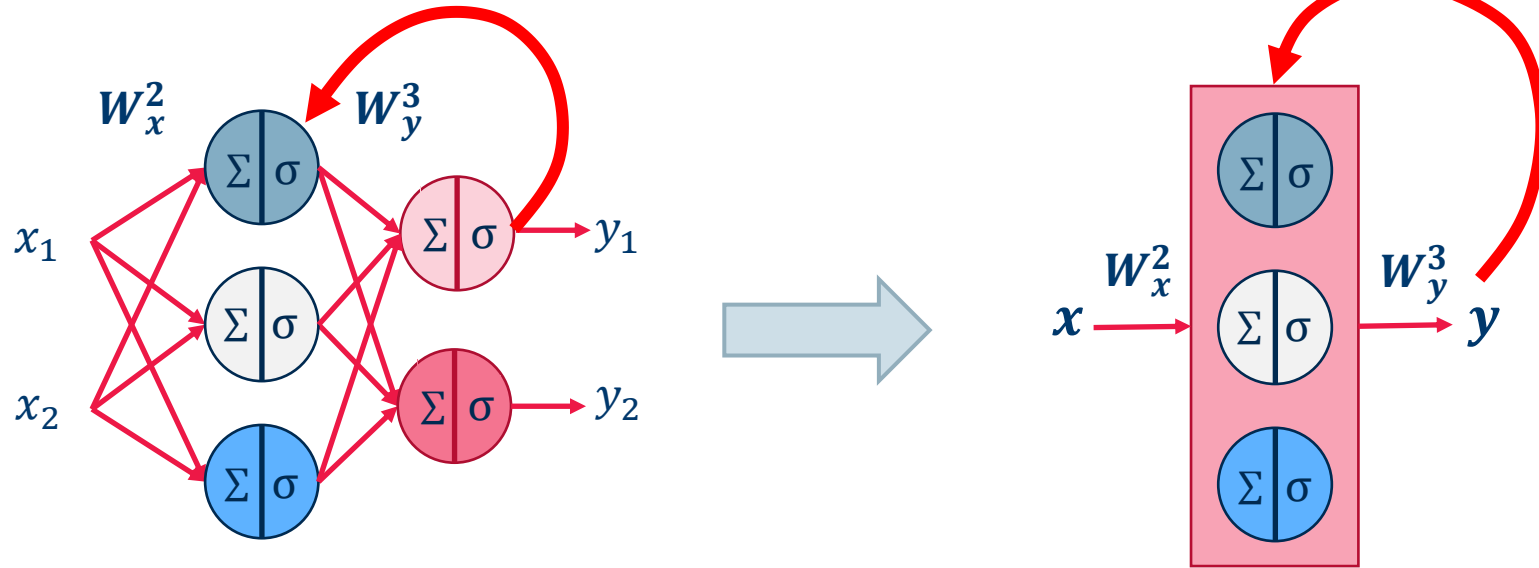


Input	Output y
Mag	like
Ich	I
Schokolade	chocolate

## From Feed Forward to Recurrent Neural Networks

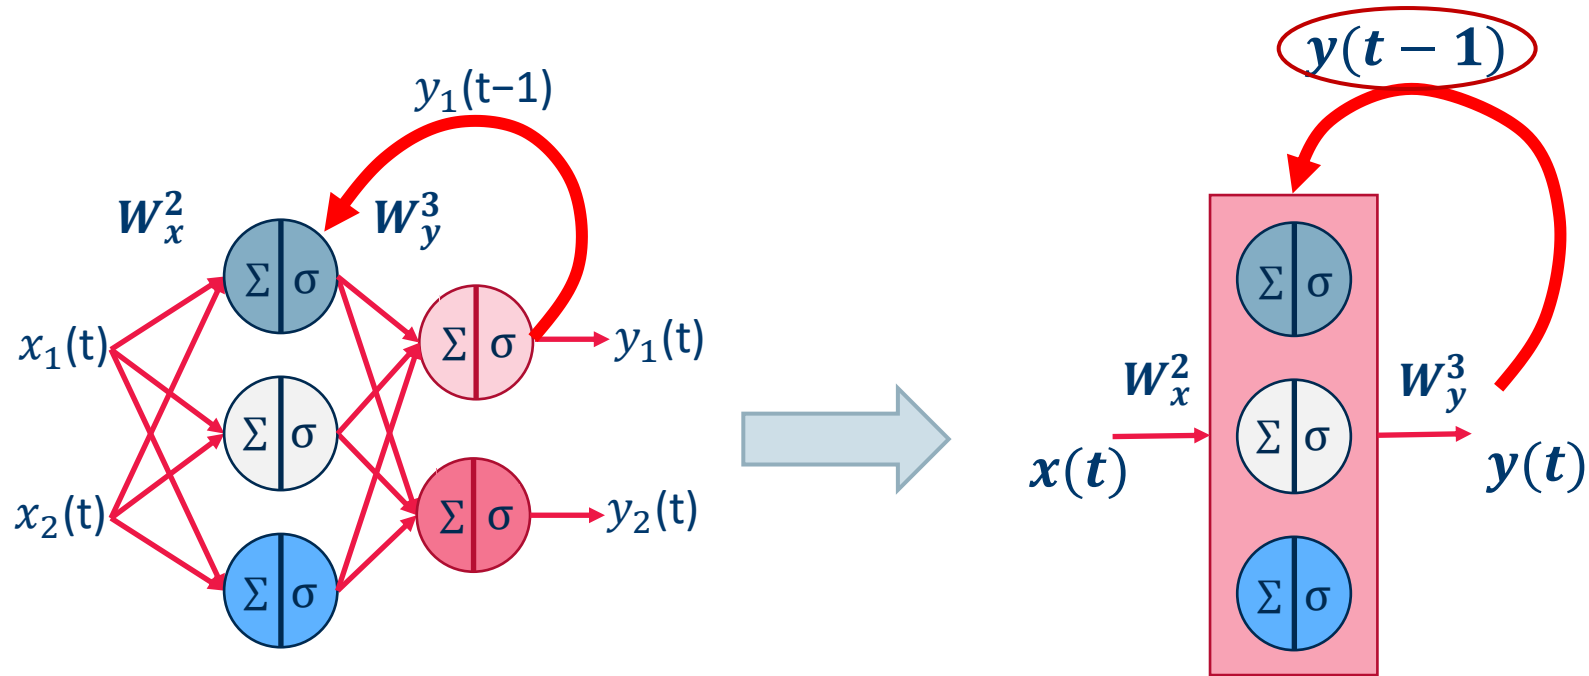


## From Feed Forward to Recurrent Neural Networks



- A Recurrent Neural Network is a FFNN with auto and/or backward connections
- Recurrent connections introduce the concept of **time** in FFNNs

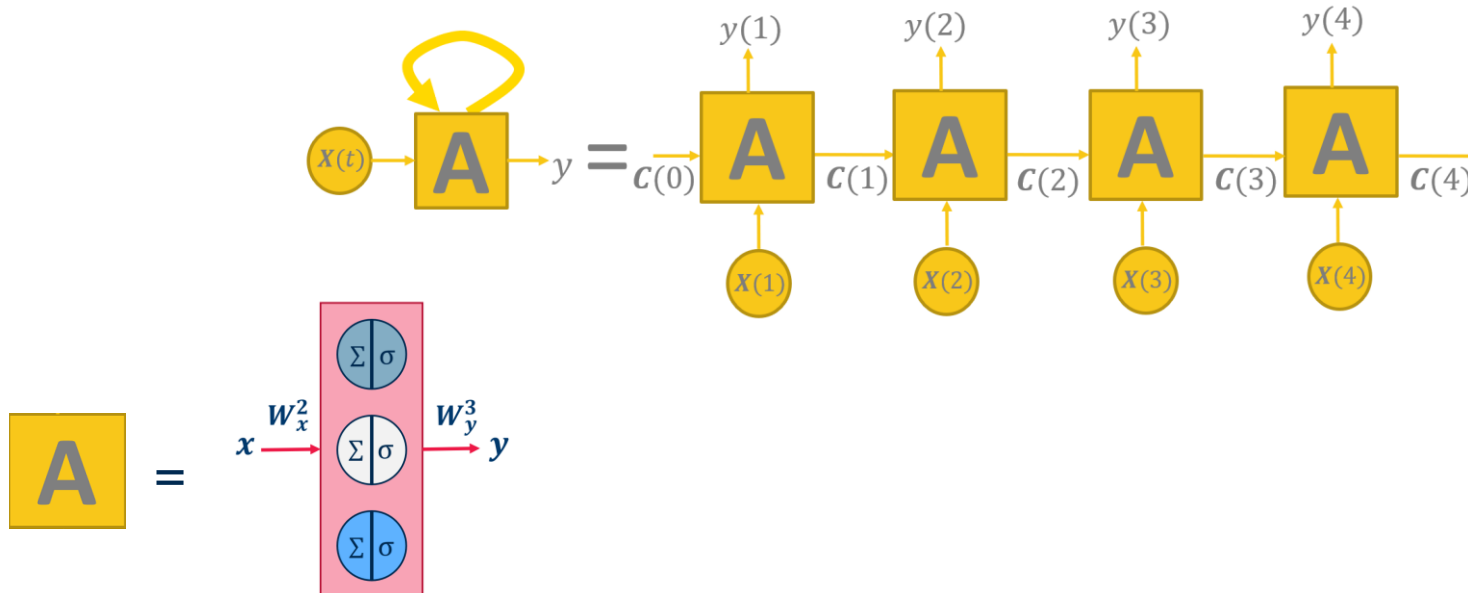
## From Feed Forward to Recurrent Neural Networks



- A Recurrent Neural Network is a FFNN with auto and/or backward connections
- Recurrent connections introduce the concept of **time** in FFNNs

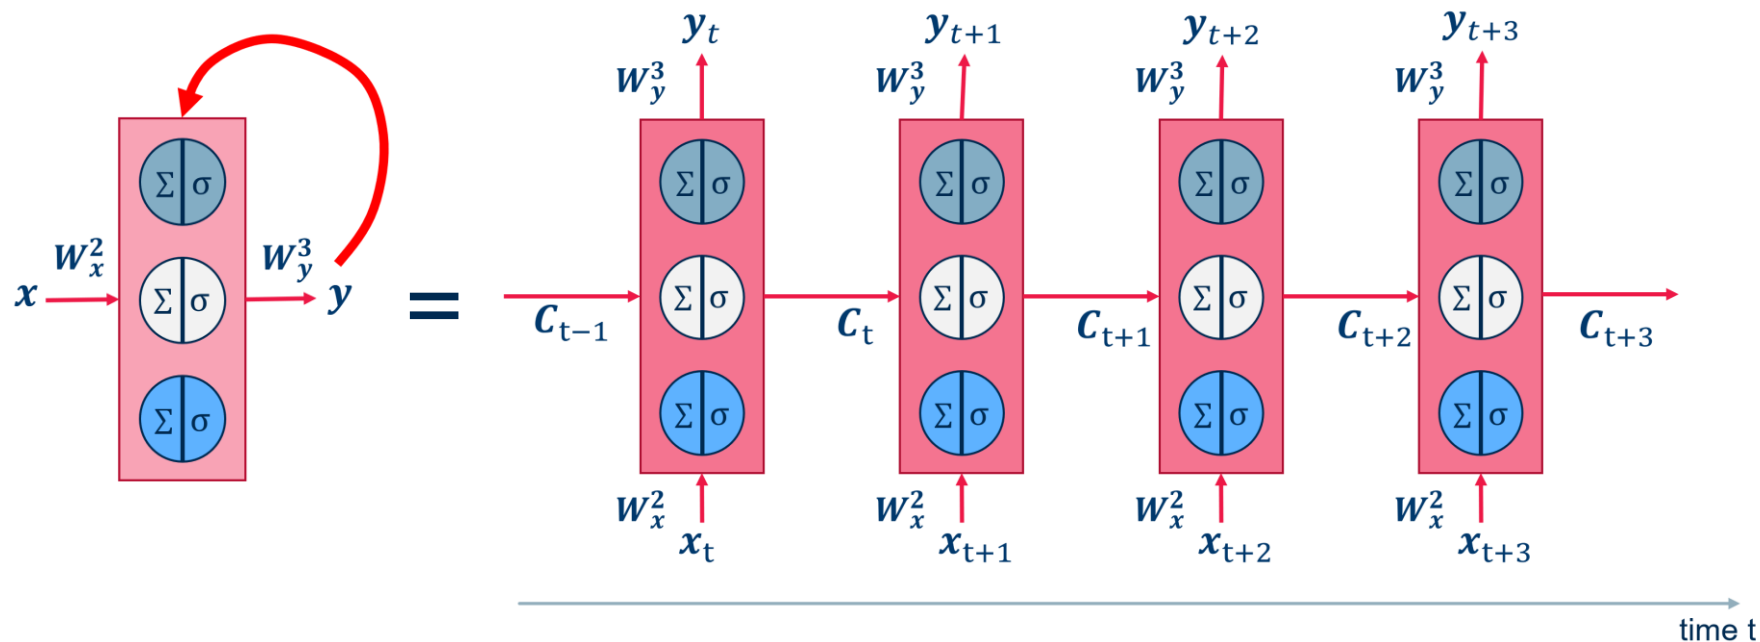
## How can we represent a RNN over time?

- At every time  $t$ , FFNN  $A$  has two inputs:
  - $\mathbf{x}(t)$
  - some shape of  $\mathbf{y}(t-1)$  -> state of network  $A$ :  $\mathbf{C}(t-1)$
- The recurrent network can then be **unrolled** over time around  $A$



## Unrolling of a RNN over time

The unrolled version of the original network in  $m$  intermediate steps becomes a FFNN and can be trained with BackPropagation: **Back-Propagation Through Time (BPTT)**.



## Summarizing: RNNs and BPTT

- Neural network architectures with recurring connections on some units are named Recurrent Neural Networks (RNNs).
- Adding a recurrent connection to one unit might store information about past inputs in the evolving status of the unit.
- Training set:  $\{\mathbf{X}(t), \mathbf{y}(t)\}$  for  $t=1, 2, \dots, N$ .
- For each  $\mathbf{X}(t)$ , the recurrent connection requires  $m$  steps into the future to produce the final output.
- An easy trick to represent the recurrent network is to unroll it into  $m$  copies of the feedforward internal block “A”, each with their set of static weight matrix  $\mathbf{W}$ . Each copy of “A” receives inputs  $\mathbf{X}(t)$  and  $\mathbf{C}(t-1)$  and produces output  $\mathbf{y}(t)$ .
- A modified version of the Back-Propagation algorithm is used to train the unrolled version in  $m$  intermediate steps of the original neural architecture: Back-Propagation Through Time (BPTT).

# Long Short Term Memory



# Simple Recurrent Unit

The simplest possible recurrent unit is a neuron with an auto-connection.

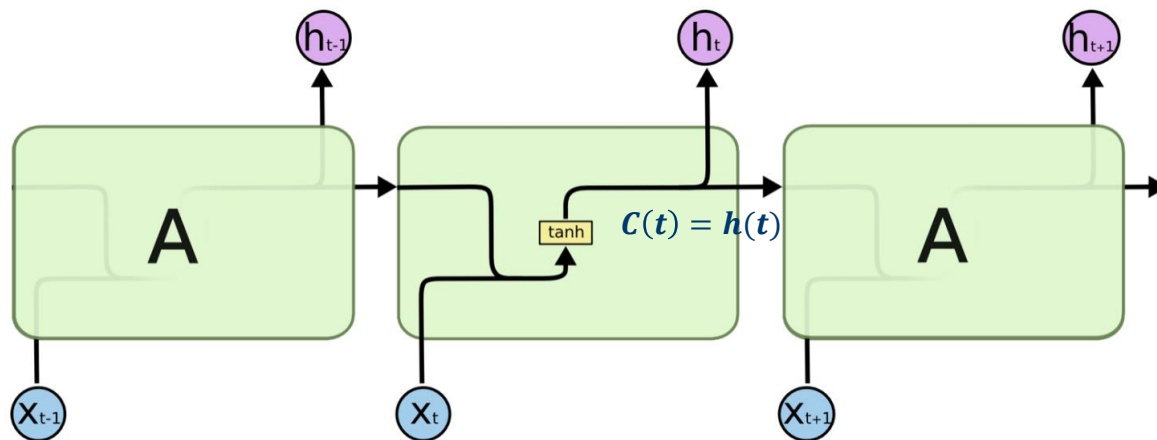
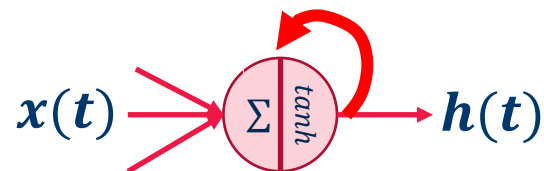


Image Source: Christopher Olah, <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

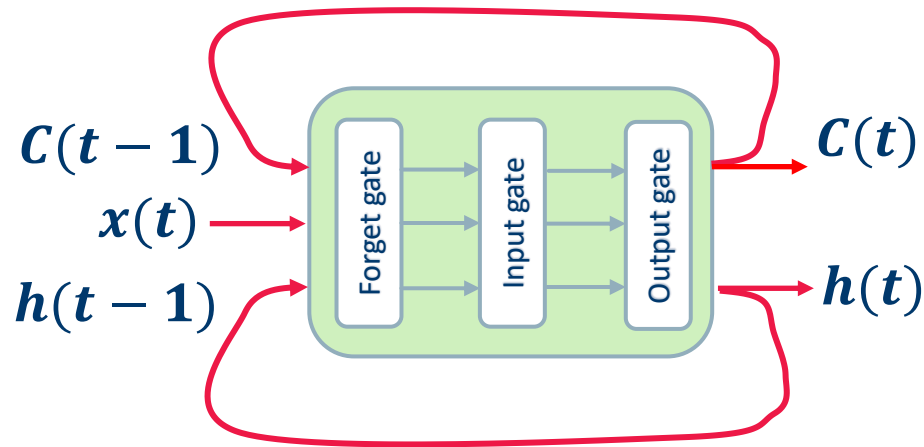
The “memory” of simple RNNs is sometimes too limited to be useful:

- *“Cars drive on the \_\_\_\_” (road)*
- *“I love the beach.  
My favorite sound is the crashing of the \_\_\_\_” (cars? glass? waves?)*
- Sometimes we need to go back deeper in time

# LSTM = Long Short Term Memory

This is an engineered type of unit with three gates:

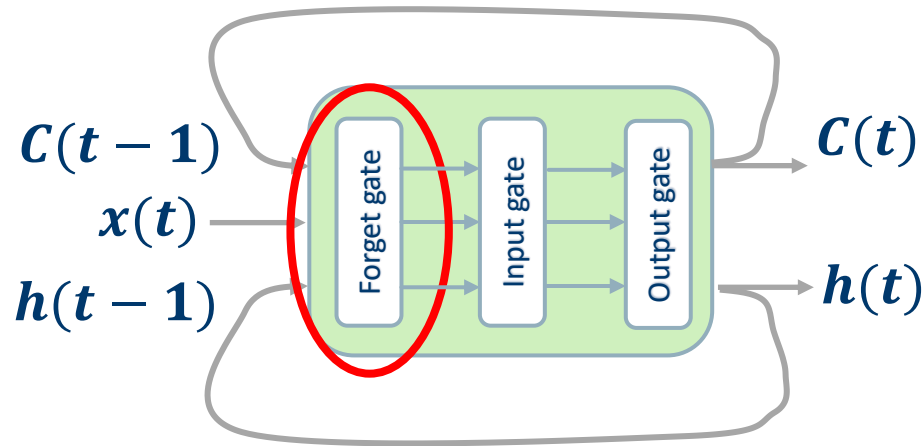
- Forget gate
- Input gate
- Output gate



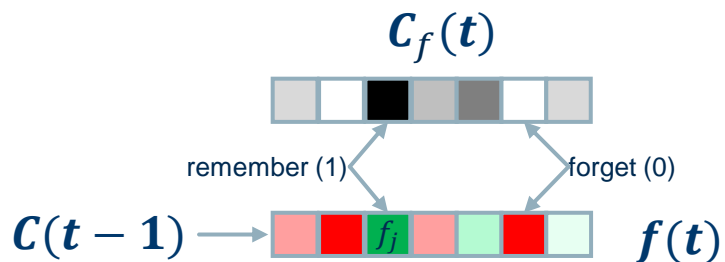
# LSTM = Forget Gate

This is an engineered type of unit with three gates:

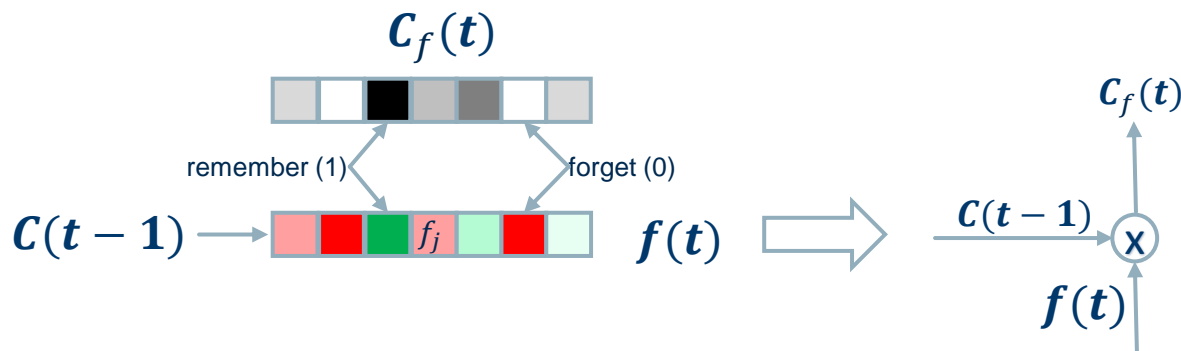
- Forget gate
- Input gate
- Output gate



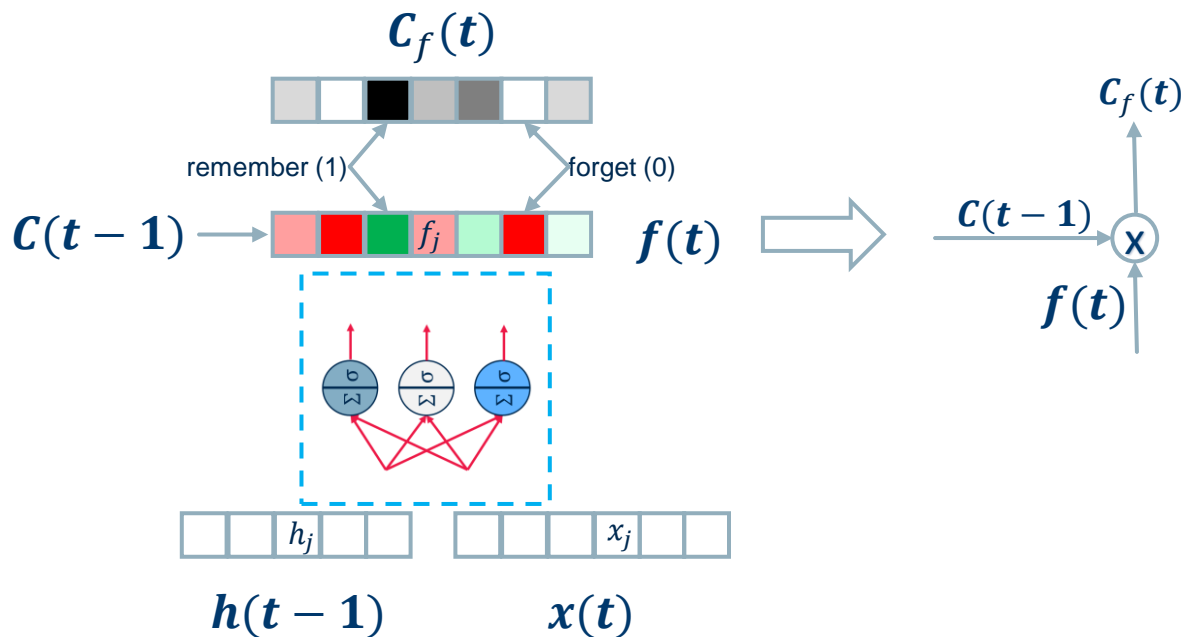
- **Forget Gate** is trained to forget the status.
- At time  $t$ , the forget gate decides which item of  $\mathcal{C}(t - 1)$  to keep (and how much of it) in  $\mathcal{C}(t)$ , given input vector  $x(t)$  and previous output  $h(t - 1)$ .



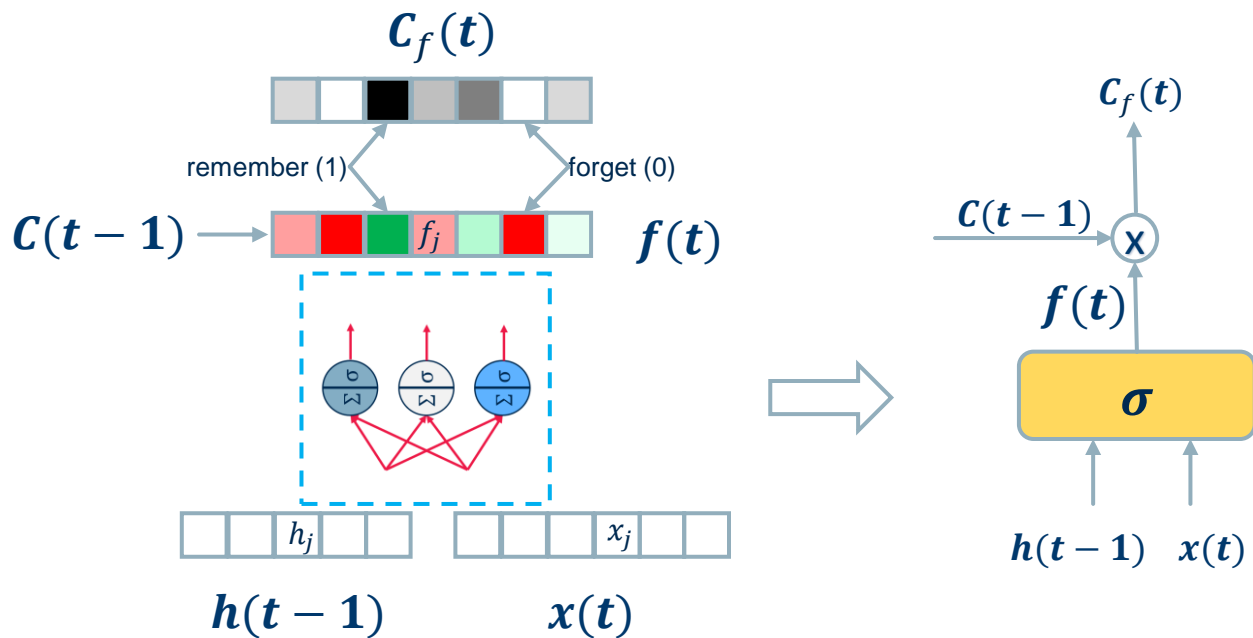
- **Forget Gate** is trained to forget the status.
- At time  $t$ , the forget gate decides which item of  $\mathcal{C}(t - 1)$  to keep (and how much of it) in  $\mathcal{C}(t)$ , given input vector  $x(t)$  and previous output  $h(t - 1)$ .



- **Forget Gate** is trained to forget the status.
- At time  $t$ , the forget gate decides which item of  $\mathcal{C}(t - 1)$  to keep (and how much of it) in  $\mathcal{C}(t)$ , given input vector  $x(t)$  and previous output  $h(t - 1)$ .

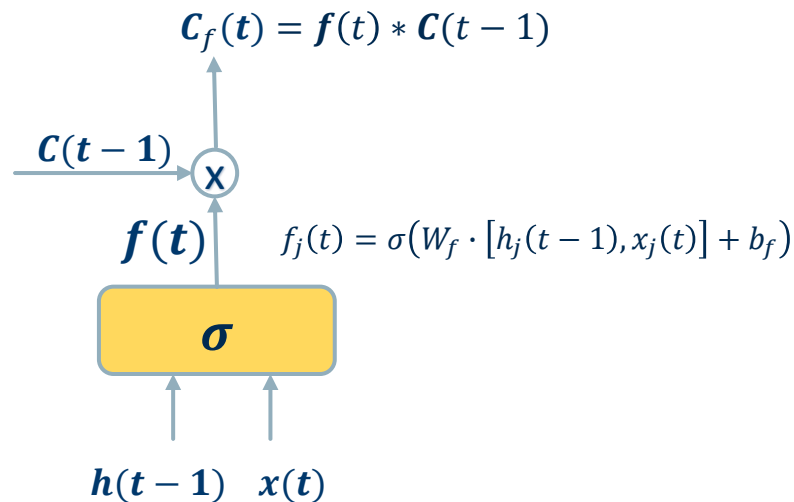
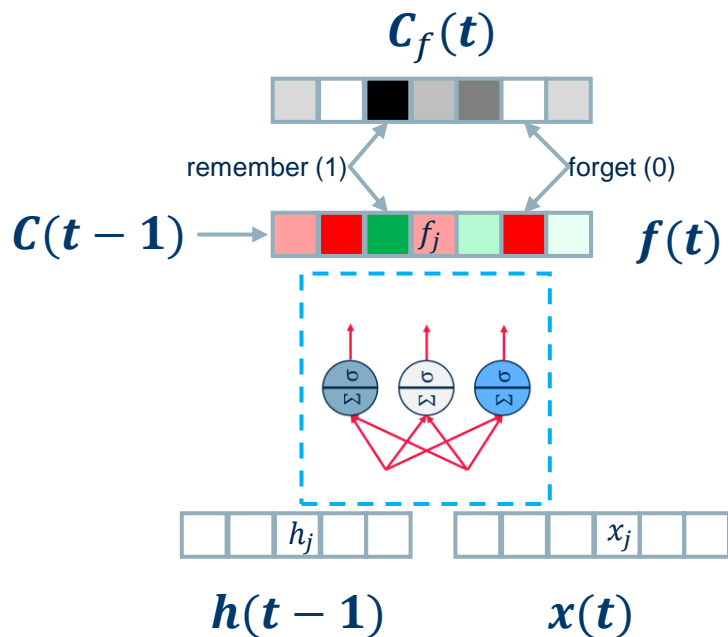


- **Forget Gate** is trained to forget the status.
- At time  $t$ , the forget gate decides which item of  $\mathcal{C}(t - 1)$  to keep (and how much of it) in  $\mathcal{C}(t)$ , given input vector  $x(t)$  and previous output  $h(t - 1)$ .





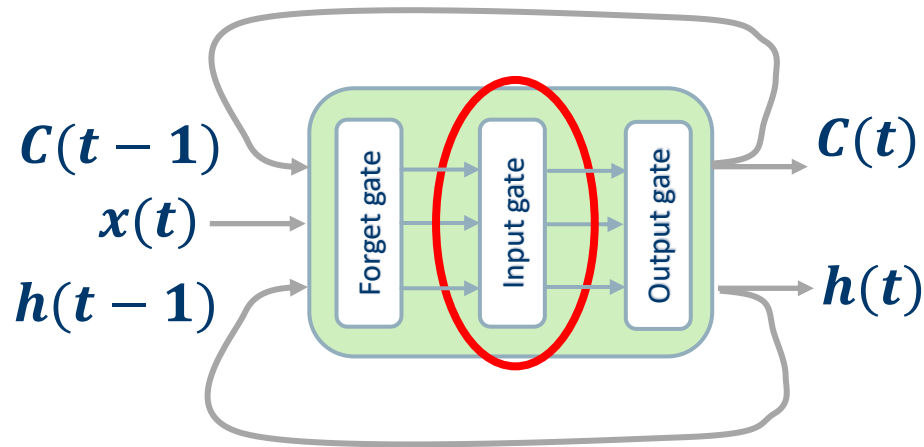
- **Forget Gate** is trained to forget the status.
- At time  $t$ , the forget gate decides which item of  $\mathcal{C}(t - 1)$  to keep (and how much of it) in  $\mathcal{C}(t)$ , given input vector  $x(t)$  and previous output  $h(t - 1)$ .



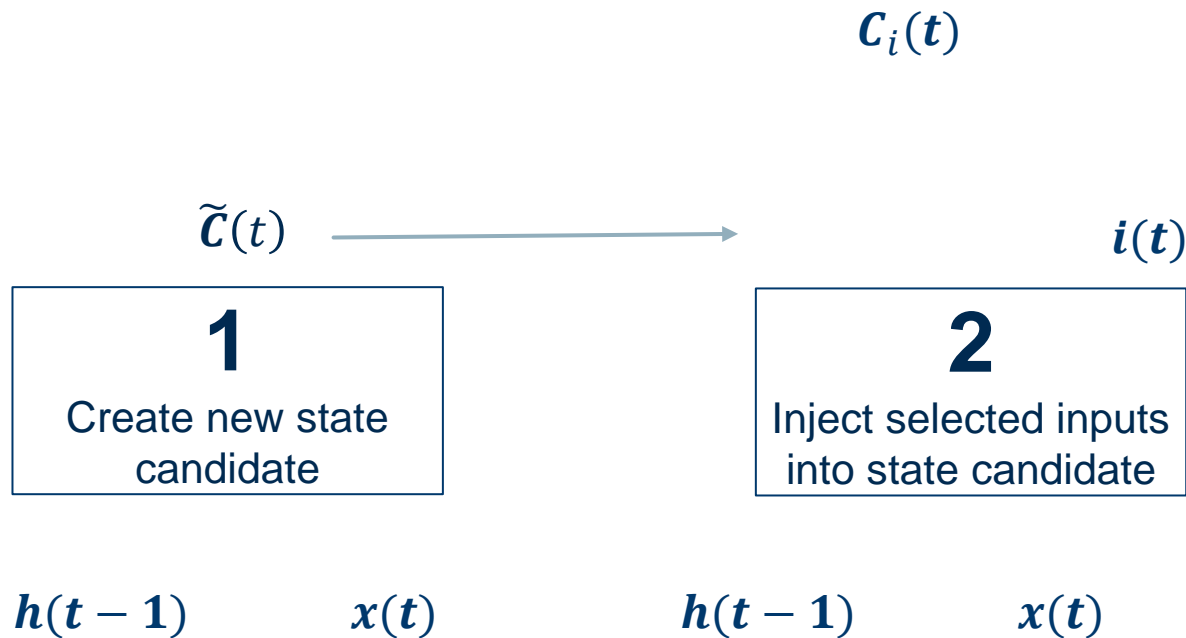
## LSTM = Input Gate

This is an engineered type of unit with three gates:

- Forget gate
- Input gate
- Output gate

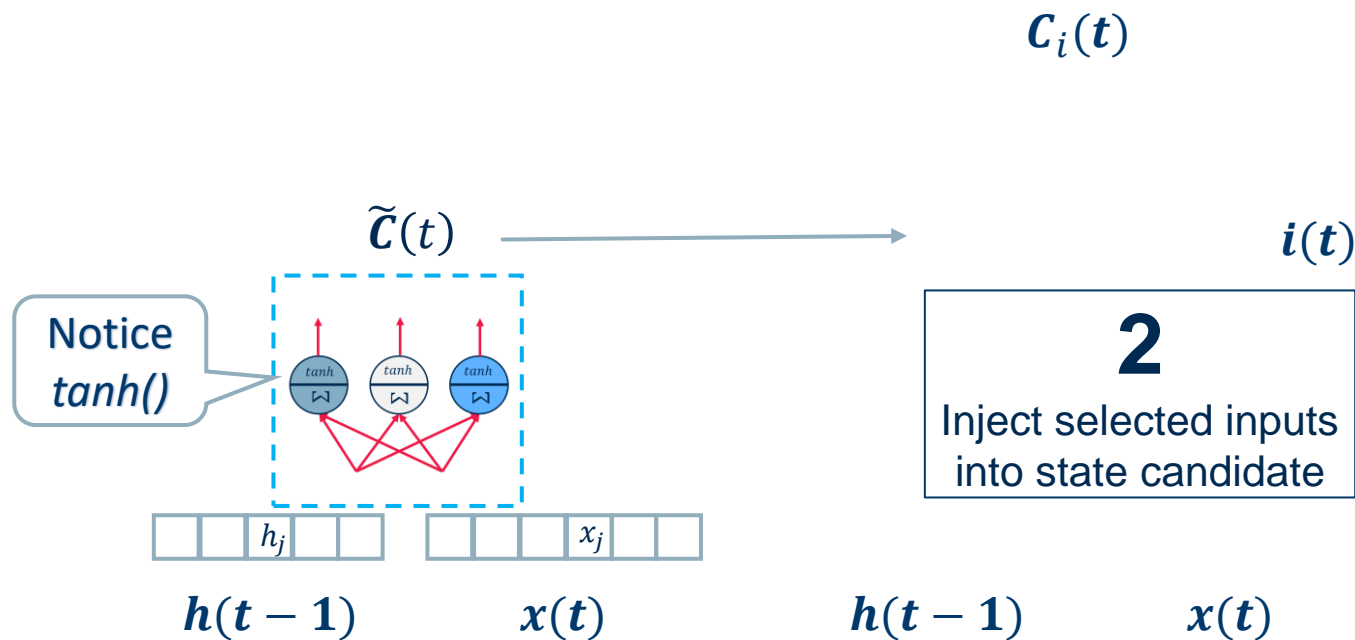


- **Input Gate** is trained to inject significant parts of the current input into the status.
- At time  $t$ , the input gate decides which item of  $x(t)$  to inject (and how much of it) into  $C(t)$ , given input vector  $x(t)$  and previous output  $h(t - 1)$ .



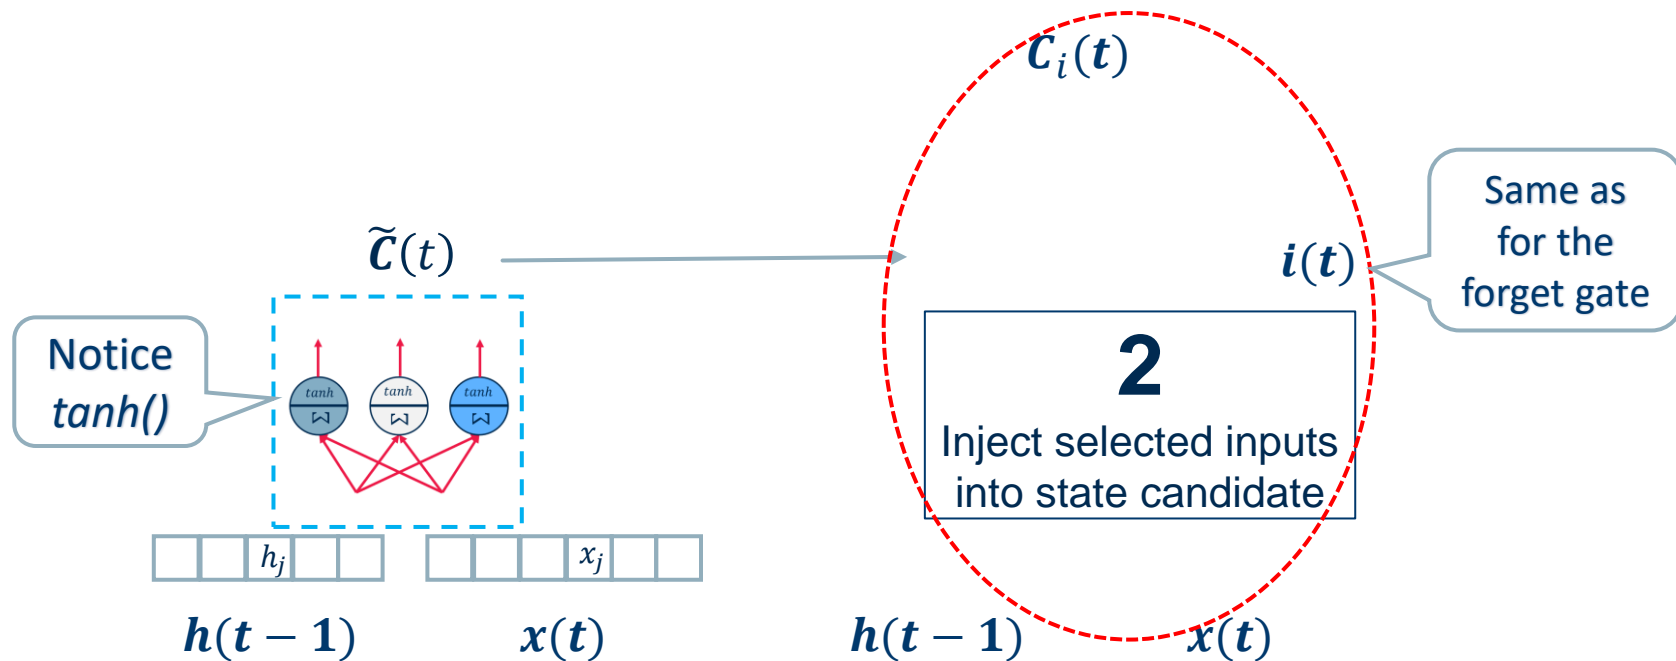
## LSTM: Input Gate – create new state candidate

- **Input Gate** is trained to inject significant parts of the current input into the status.
- At time  $t$ , the input gate decides which item of  $x(t)$  to inject (and how much of it) into  $C(t)$ , given input vector  $x(t)$  and previous output  $h(t-1)$ .



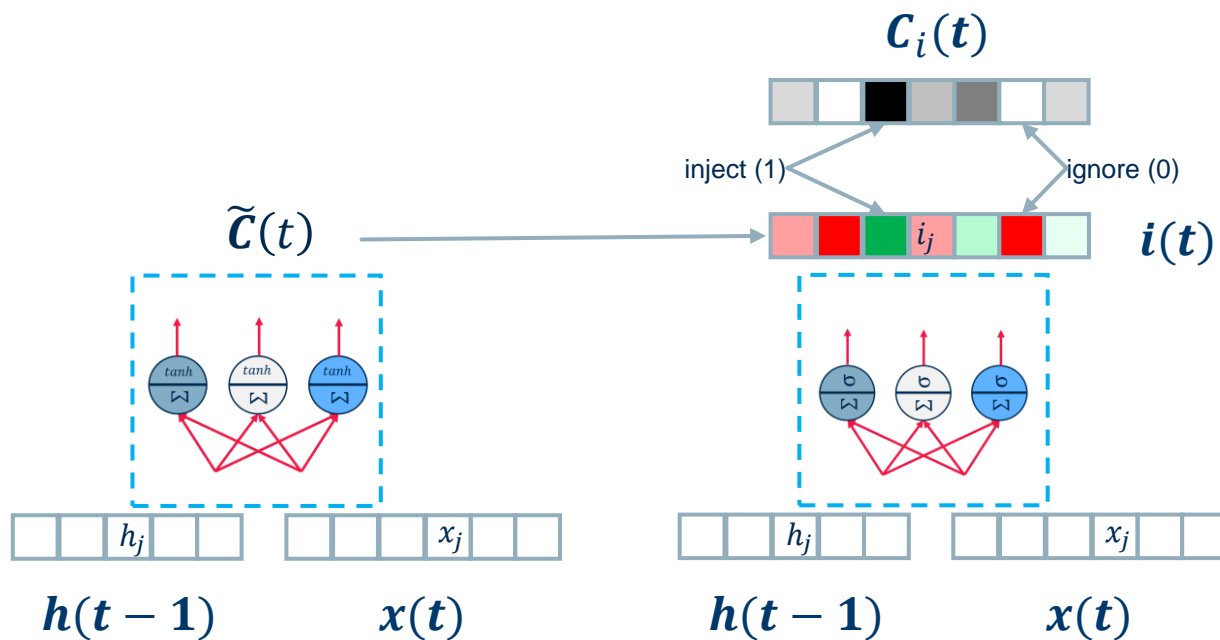
## LSTM: Input Gate – inject input

- **Input Gate** is trained to inject significant parts of the current input into the status.
- At time  $t$ , the input gate decides which item of  $x(t)$  to inject (and how much of it) into  $C(t)$ , given input vector  $x(t)$  and previous output  $h(t-1)$ .



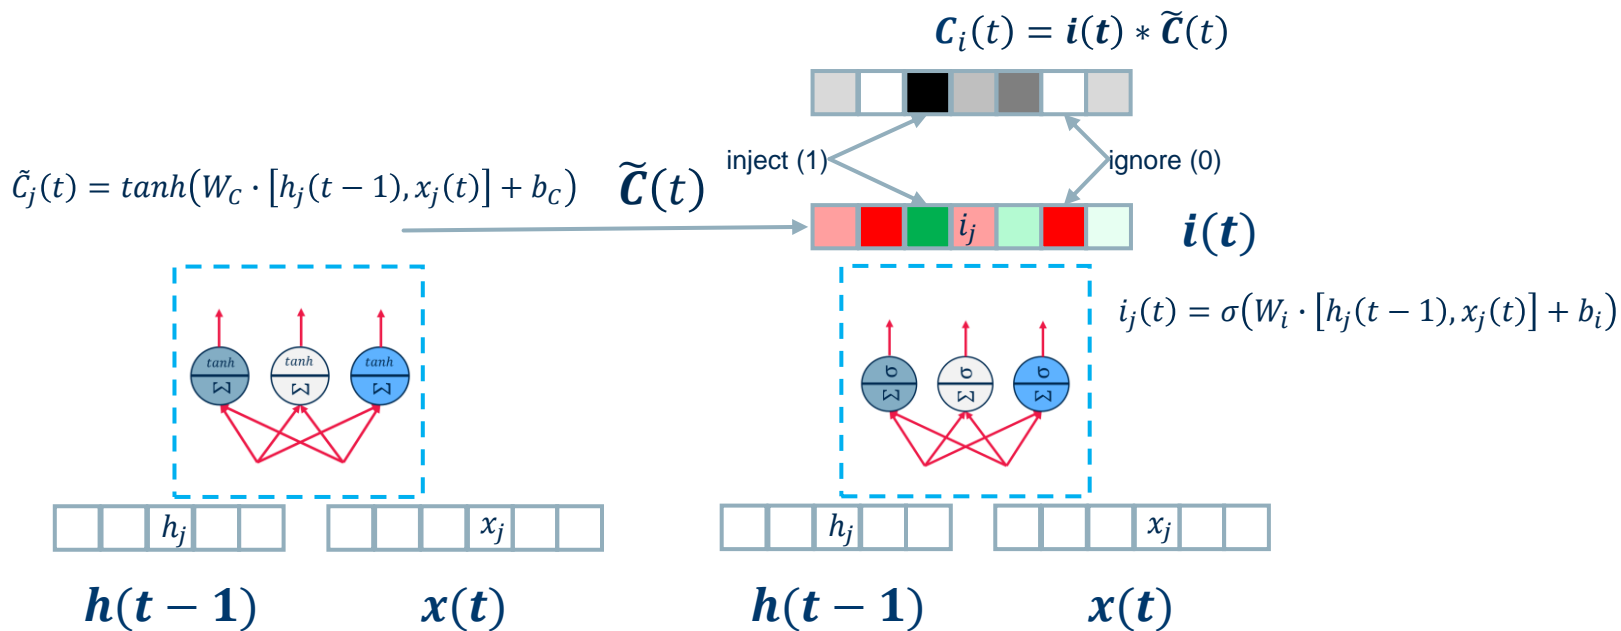
## LSTM: Input Gate – inject input

- **Input Gate** is trained to inject significant parts of the current input into the status.
- At time  $t$ , the input gate decides which item of  $x(t)$  to inject (and how much of it) into  $C(t)$ , given input vector  $x(t)$  and previous output  $h(t-1)$ .

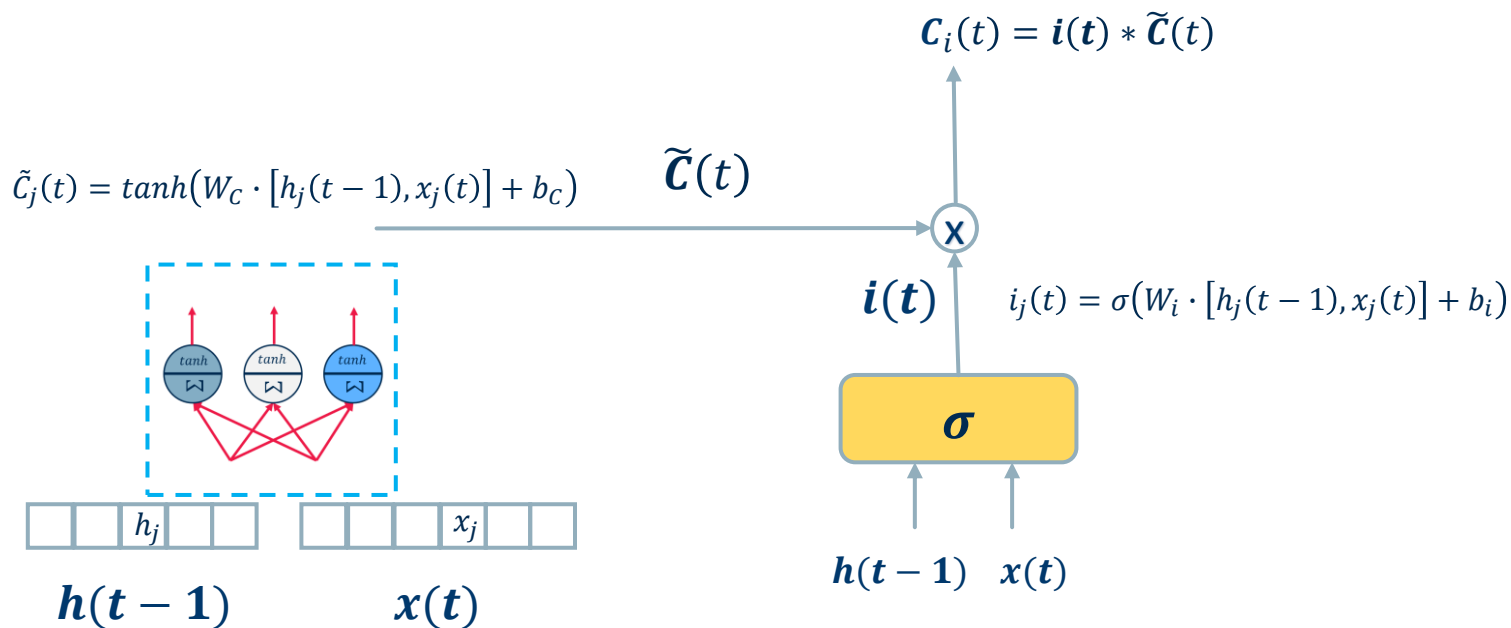


## LSTM: Input Gate

- **Input Gate** is trained to inject significant parts of the current input into the status.
- At time  $t$ , the input gate decides which item of  $x(t)$  to inject (and how much of it) into  $C(t)$ , given input vector  $x(t)$  and previous output  $h(t-1)$ .

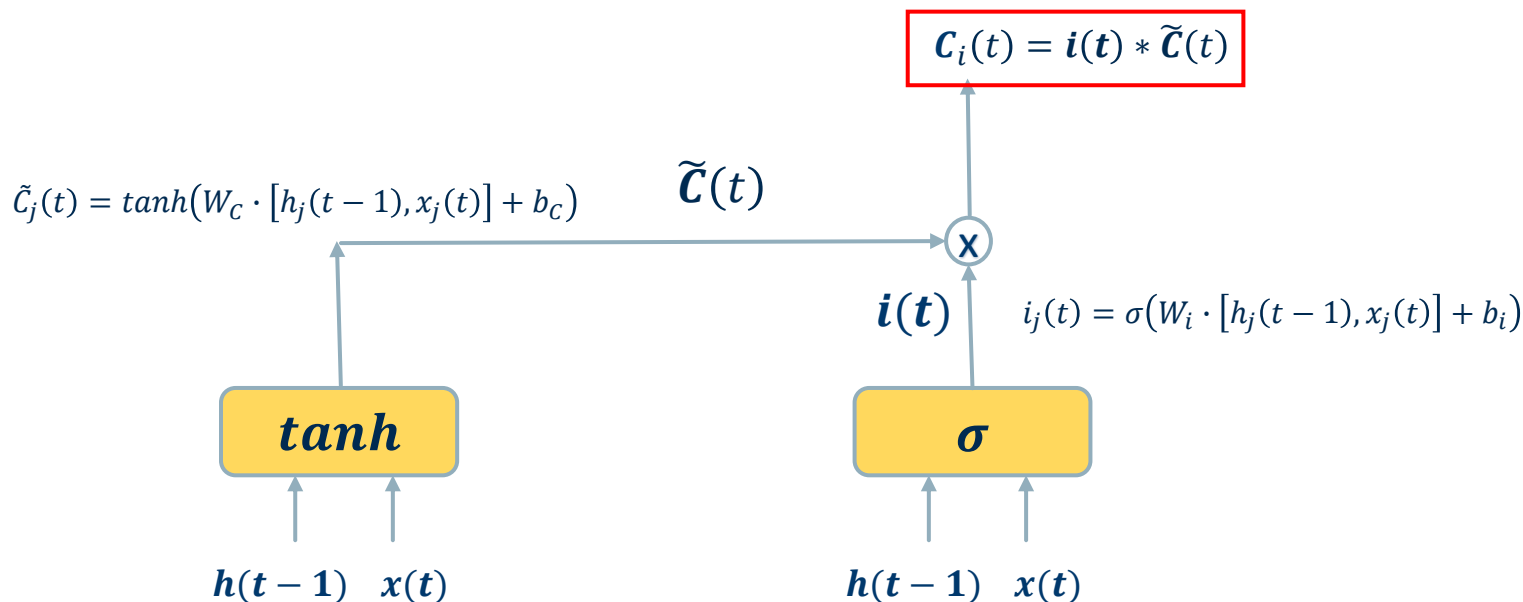


- **Input Gate** is trained to inject significant parts of the current input into the status.
- At time  $t$ , the input gate decides which item of  $x(t)$  to inject (and how much of it) into  $\mathcal{C}(t)$ , given input vector  $x(t)$  and previous output  $h(t-1)$ .





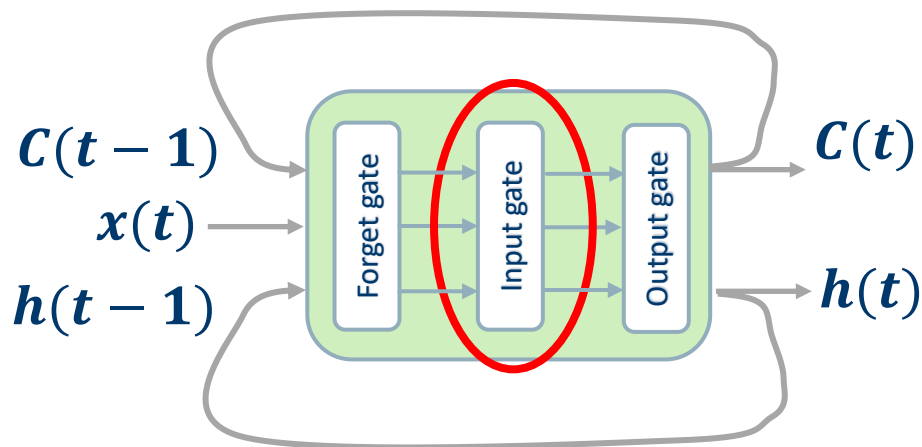
- **Input Gate** is trained to inject significant parts of the current input into the status.
- At time  $t$ , the input gate decides which item of  $x(t)$  to inject (and how much of it) into  $\mathbf{C}(t)$ , given input vector  $x(t)$  and previous output  $\mathbf{h}(t - 1)$ .



## LSTM = Input Gate

This is an engineered type of unit with three gates:

- Forget gate
- Input gate
- Output gate

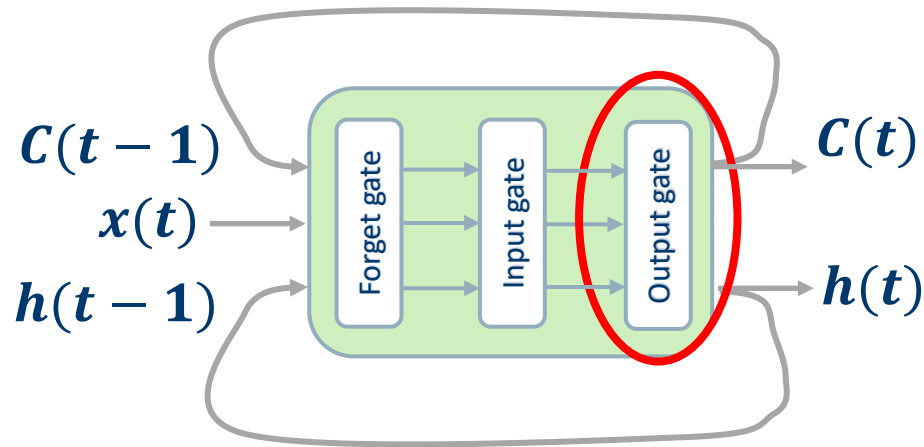


$$C(t) = C_f(t) + C_i(t) = f(t) * C(t-1) + i(t) * \tilde{C}(t)$$

# LSTM = Output Gate

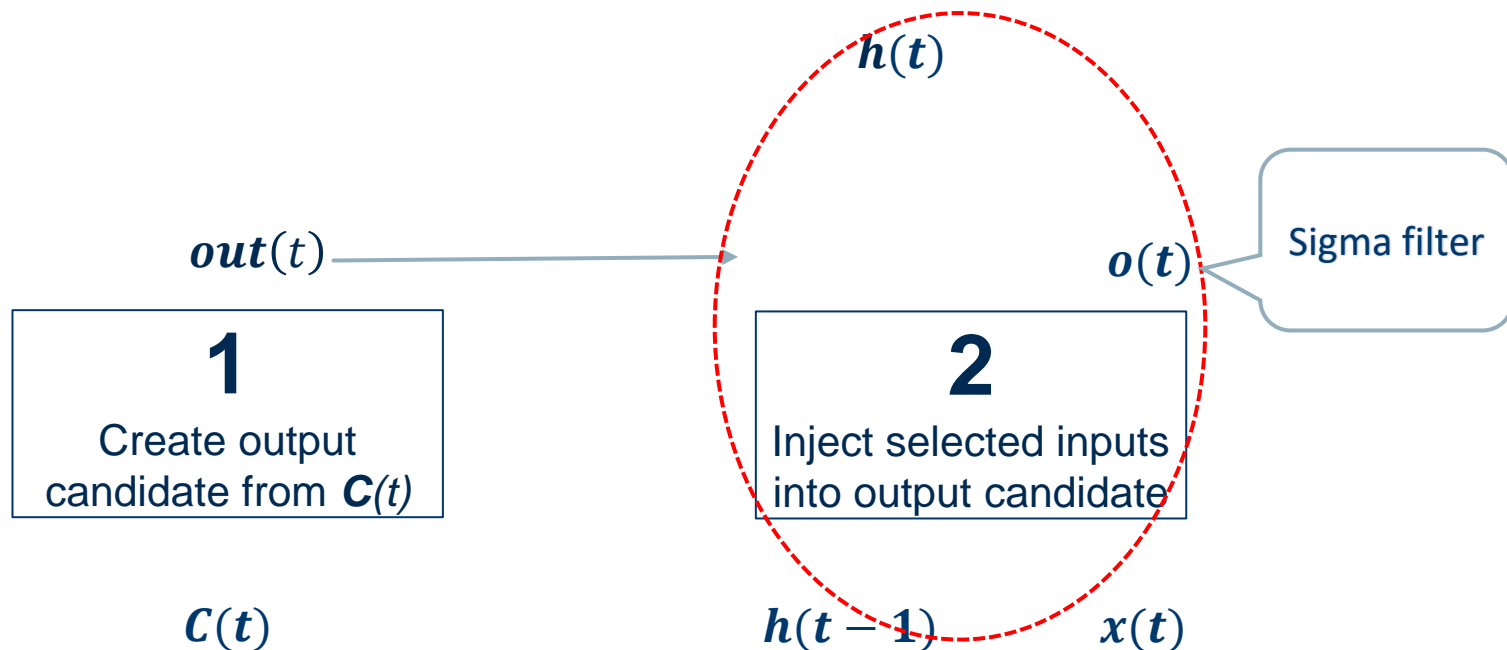
This is an engineered type of unit with three gates:

- Forget gate
- Input gate
- Output gate

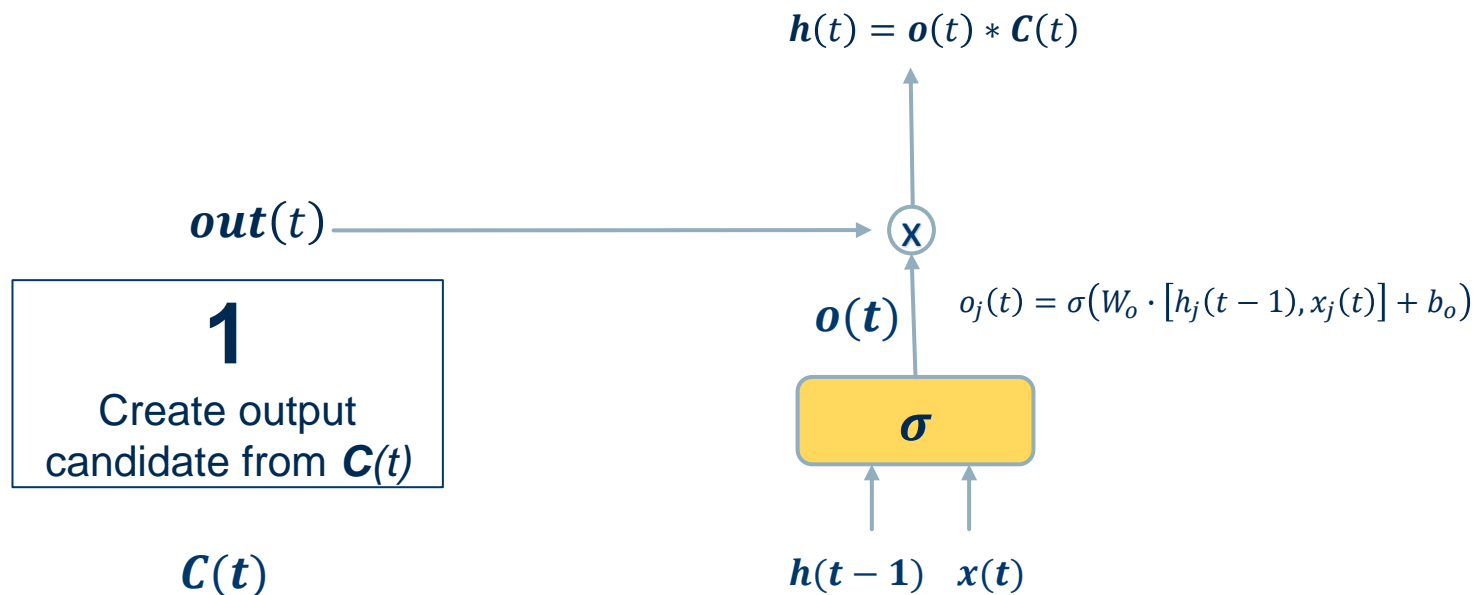


## LSTM: Output Gate – input inject into status

- **Output Gate** is trained to output a reasonable result.
- At time  $t$ , output gate decides which parts of status  $\mathbf{C}(t)$  (and how much of it) will be output, given input vector  $\mathbf{x}(t)$  and previous output  $\mathbf{h}(t - 1)$ .

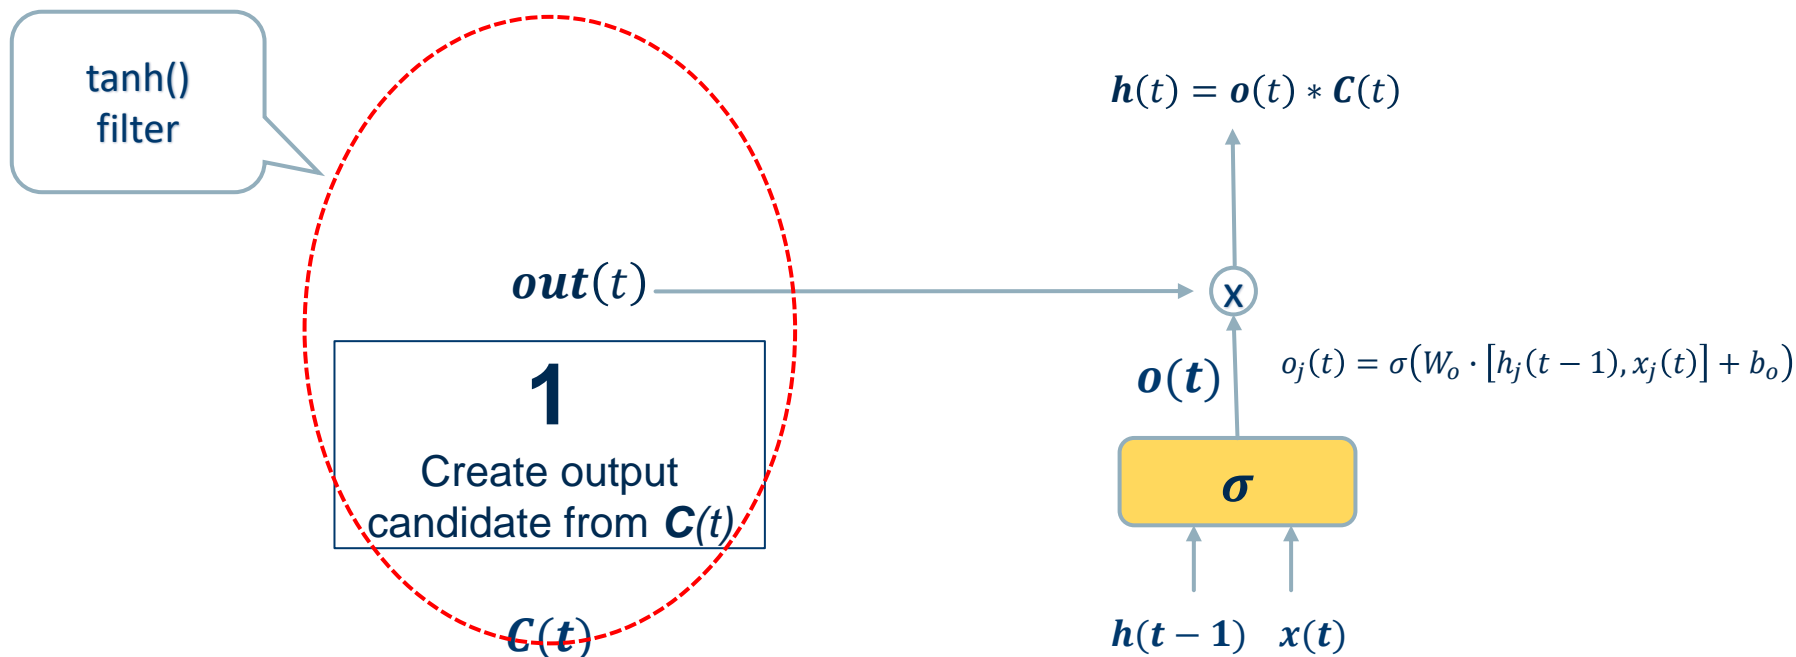


- **Output Gate** is trained to output a reasonable result.
- At time  $t$ , output gate decides which parts of status  $\mathbf{C}(t)$  (and how much of it) will be output, given input vector  $\mathbf{x}(t)$  and previous output  $\mathbf{h}(t - 1)$ .



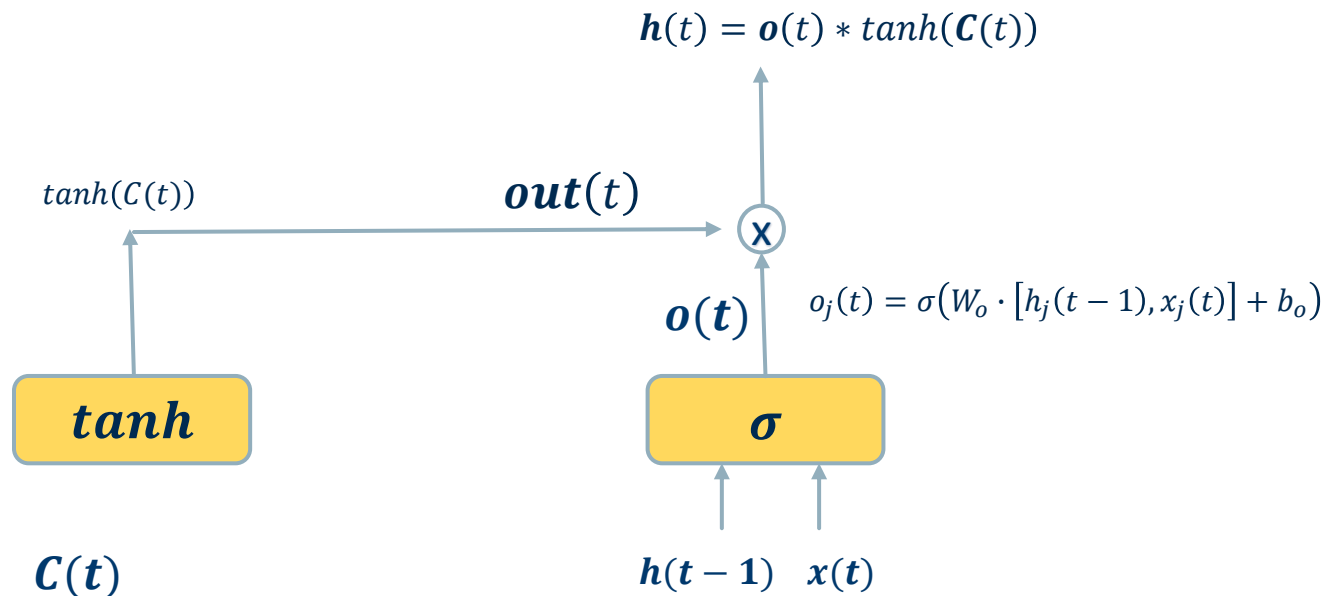
## LSTM: Input Gate – prepare output candidate

- **Output Gate** is trained to output a reasonable result.
- At time  $t$ , output gate decides which parts of status  $\mathbf{C}(t)$  (and how much of it) will be output, given input vector  $\mathbf{x}(t)$  and previous output  $\mathbf{h}(t - 1)$ .



## LSTM: Input Gate – prepare output candidate

- **Output Gate** is trained to output a reasonable result.
- At time  $t$ , output gate decides which parts of status  $\mathcal{C}(t)$  (and how much of it) will be output, given input vector  $x(t)$  and previous output  $h(t - 1)$ .



# LSTM = Long Short Term Memory

## Special type of unit with three gates

- Forget gate
- Input gate
- Output gate

Forget gate

Input gate

Output gate

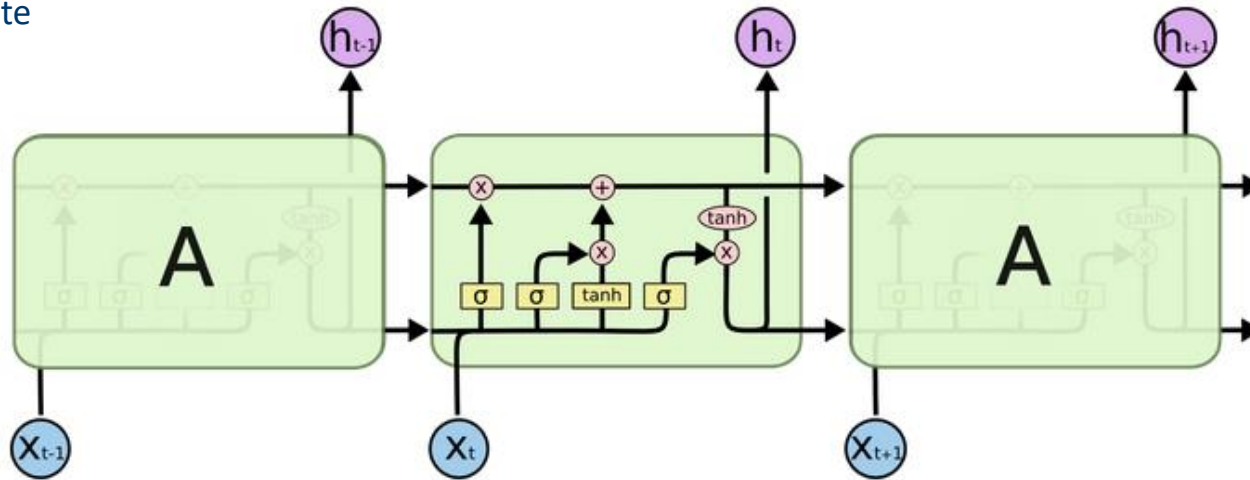
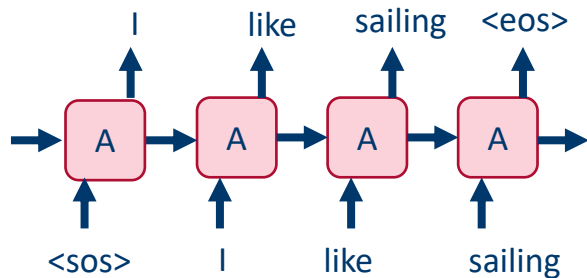


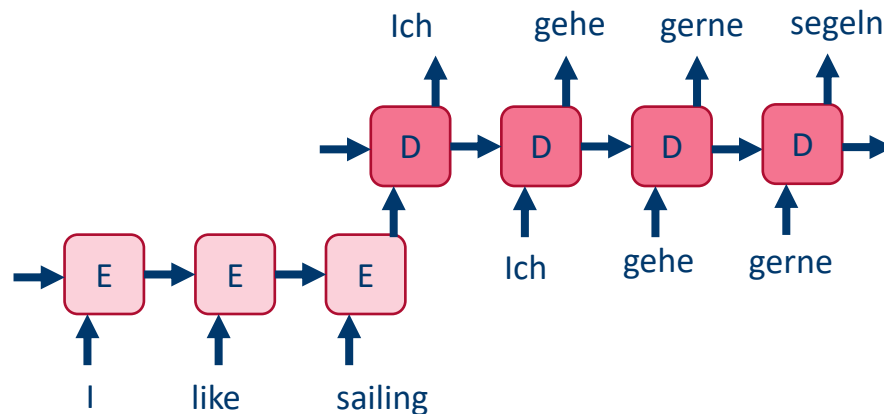
Image Source: Christopher Olah, <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>



## Many to Many



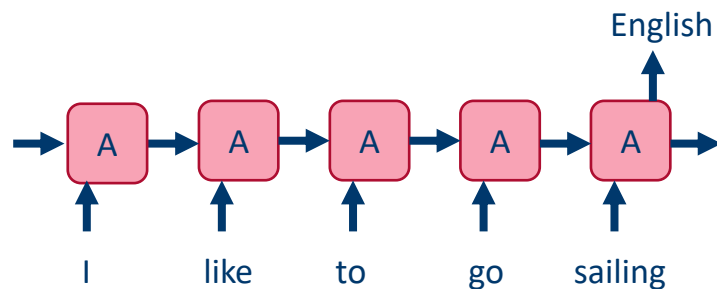
Language model



Neural machine translation

# Different Network-Structures and Applications

Many to one



Language classification  
Text classification

One to many

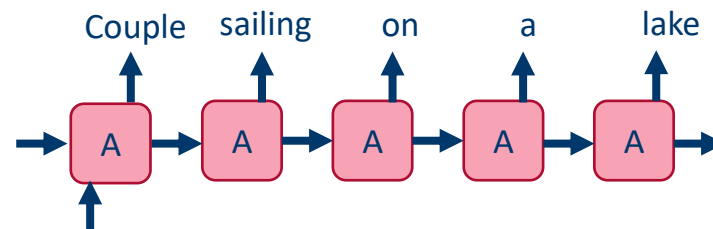


Image captioning

# Neural Network: Code-free Example

## Define Network

Keras Input Layer



Input Shape  
?, Dictionary Size

Keras LSTM Layer



Output: Sequence of  
Hidden States

Keras Dropout Layer



Regularization

Keras Dense Layer



Activation:  
Linear

Keras Dense Layer



Activation:  
Softmax

## Read and Pre-Process Input Data

Pre-Processing



## Train Network

Keras Network  
Learner



## Edit and Save Networks

DL Python  
Network Editor



Add Temperature  
and Remove Dropout

Keras to TensorFlow  
Network Converter



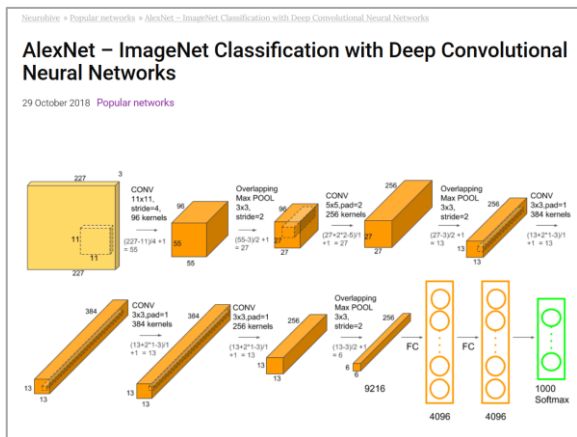
TensorFlow  
Network Writer



# Convolutional Neural Networks (CNNs)

# AlexNet & friends

- The big breakthrough in deep learning happened in 2012 with deep convolutional neural networks
- Here deep learning based AlexNet network won the ImageNet challenge with an unprecedented margin.
- The top-five error rate of AlexNet was 15 percent, while the next best competitor ended up with 26 percent.
- This victory kicked off the surge in deep learning networks.



<https://neurohive.io/en/popular-networks/alexnet-imagenet-classification-with-deep-convolutional-neural-networks/>

## Convolutional Neural Networks - CNN

- Inspired by the organization of the visual cortex in the human brain, convolutional layers simulate the concept of a receptive field.
- Individual neurons in the convolutional layer respond only when a specific area of the image (the visual field) is active.
- An array of such neurons covers the entire image by responding to slightly overlapping separated areas of the input image.

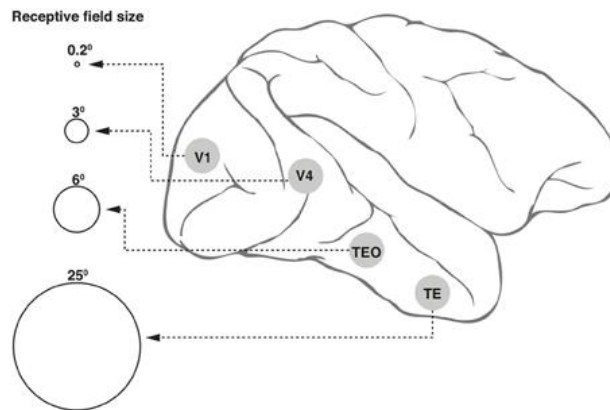
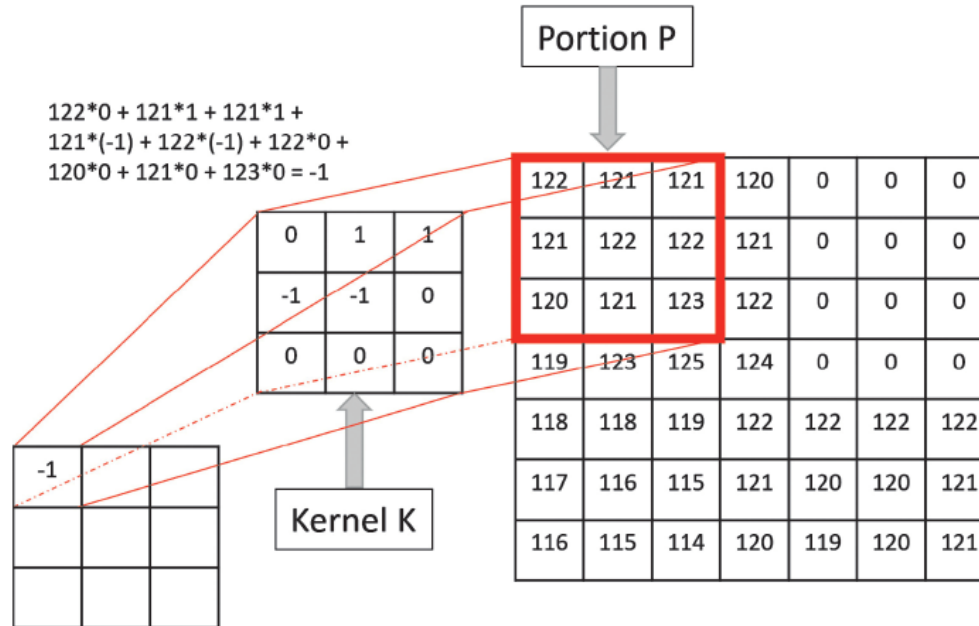


Image from: Wikimedia commons –

[https://commons.wikimedia.org/wiki/File:Receptive\\_field\\_sizes\\_along\\_the\\_ventral\\_cortical\\_stream\\_in\\_the\\_primate.jpg](https://commons.wikimedia.org/wiki/File:Receptive_field_sizes_along_the_ventral_cortical_stream_in_the_primate.jpg)

- The idea of convolution relies on a kernel  $K$ , a mask to overlap onto a portion  $P$  of the image pixels for the convolution operation.
- From the product of the kernel  $K$  and the pixels in portion  $P$  we get a number, which will be the output of the first neuron in the convolutional layer.
- Then the kernel  $K$  moves  $n$  steps on the right and goes to cover another portion  $P$  of the image possibly slightly overlapping with the previous one; the output for the second unit of the convolutional layer is generated.
- And so on till the whole image has been covered by the kernel  $K$  and convoluted into output values.
- The distance in number of pixels  $n$  between two adjacent portions  $P$  is called *stride*.

## Convolutional neurons: Example





# Convolutional Neural Networks (CNN)

- Used when data has spatial relationships, e.g. images
- Instead of connecting every neuron to the new layer a sliding window is used
- Some convolutions may detect edges or corners, while others may detect cats, dogs, or street signs inside an image

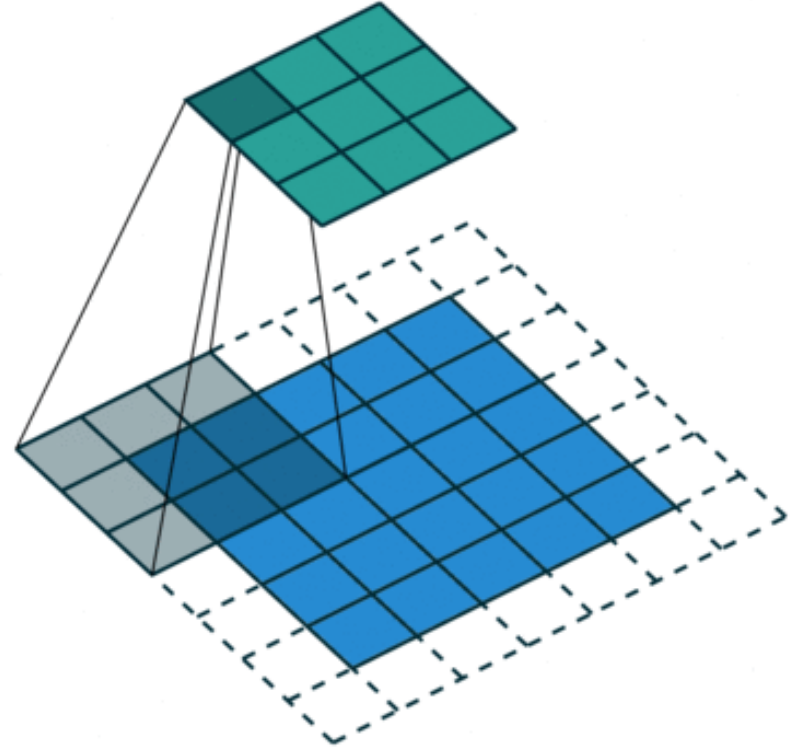


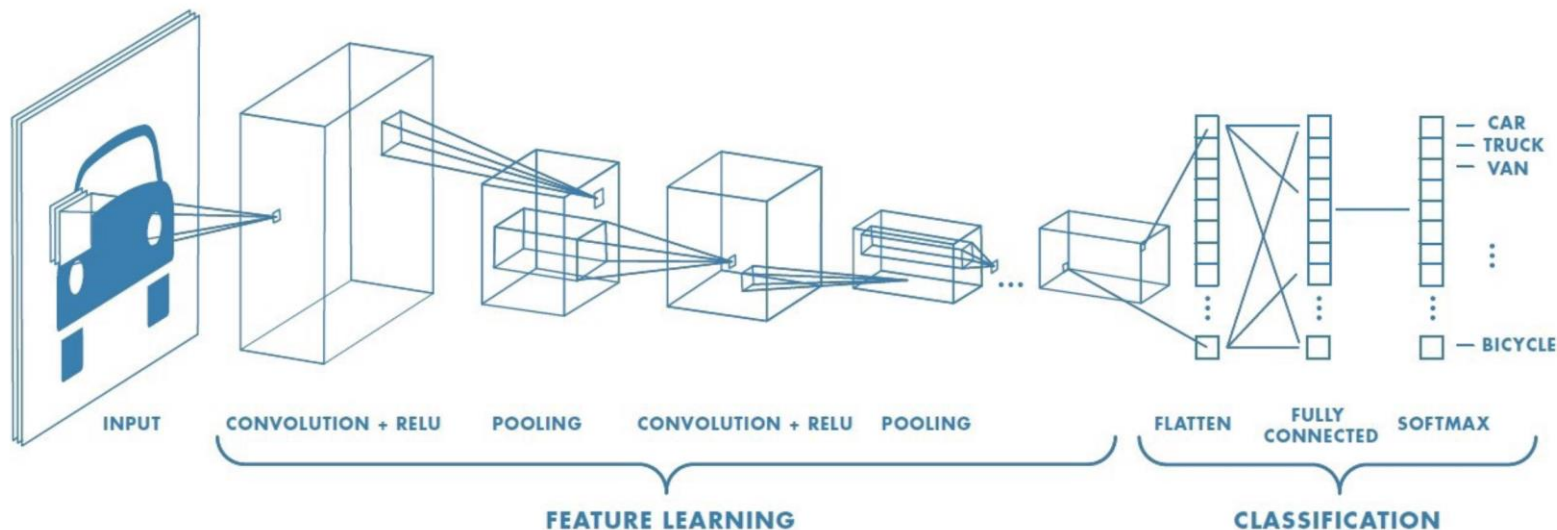
Image from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

## Pooling Layers

- Usually a number of convolutional layers are used.
- Each layer provides one further step in the process of extracting high-level features from the input image (colors, edges, entities, ...).
- After each convolutional layer, a *pooling layer* is often applied to reduce even further the data dimensionality.
- Two types of Pooling
  - **Max Pooling** returns the **maximum value** from the portion of the image covered by the Kernel.
  - **Average Pooling** returns the **average of all values** from the portion of the image covered by the Kernel.

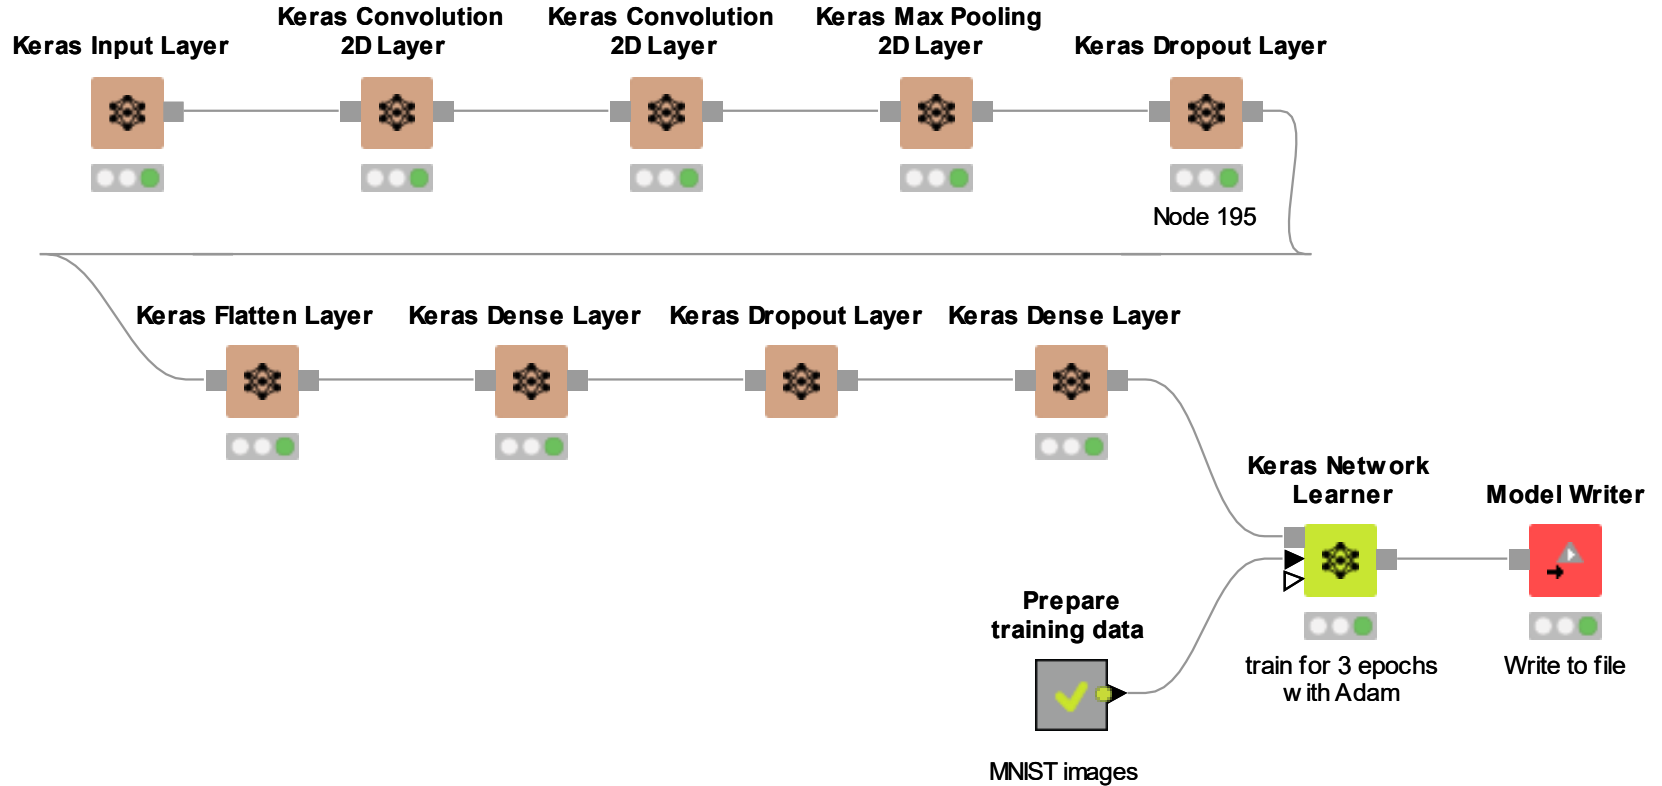
# Classification Layers

- After the sequence of convolutional + pooling layers, a classic feedforward multilayer Perceptron network is applied to carry out the classification process.
- Successful examples of CNNs for image recognition : LeNet, AlexNet, VGGNet, GoogLeNet, ResNet, ZFNet.



- Training such networks is a long and complex process, requiring very powerful machines.
- Instead of retraining a new network completely from scratch, we could recycle existing networks, already built and trained by others on **similar** data.
- This technique is called ***Transfer Learning***.
- In Transfer Learning a model developed for a task is reused as the starting point for another model on a second task.
- On top of a previously trained network we add one or more neural layers
- We freeze all or some of the previously trained layers
- And we retrain only the remaining part of the whole network on our new task

# Building CNNs with KNIME



# Generative-Adversarial Networks (GANs)

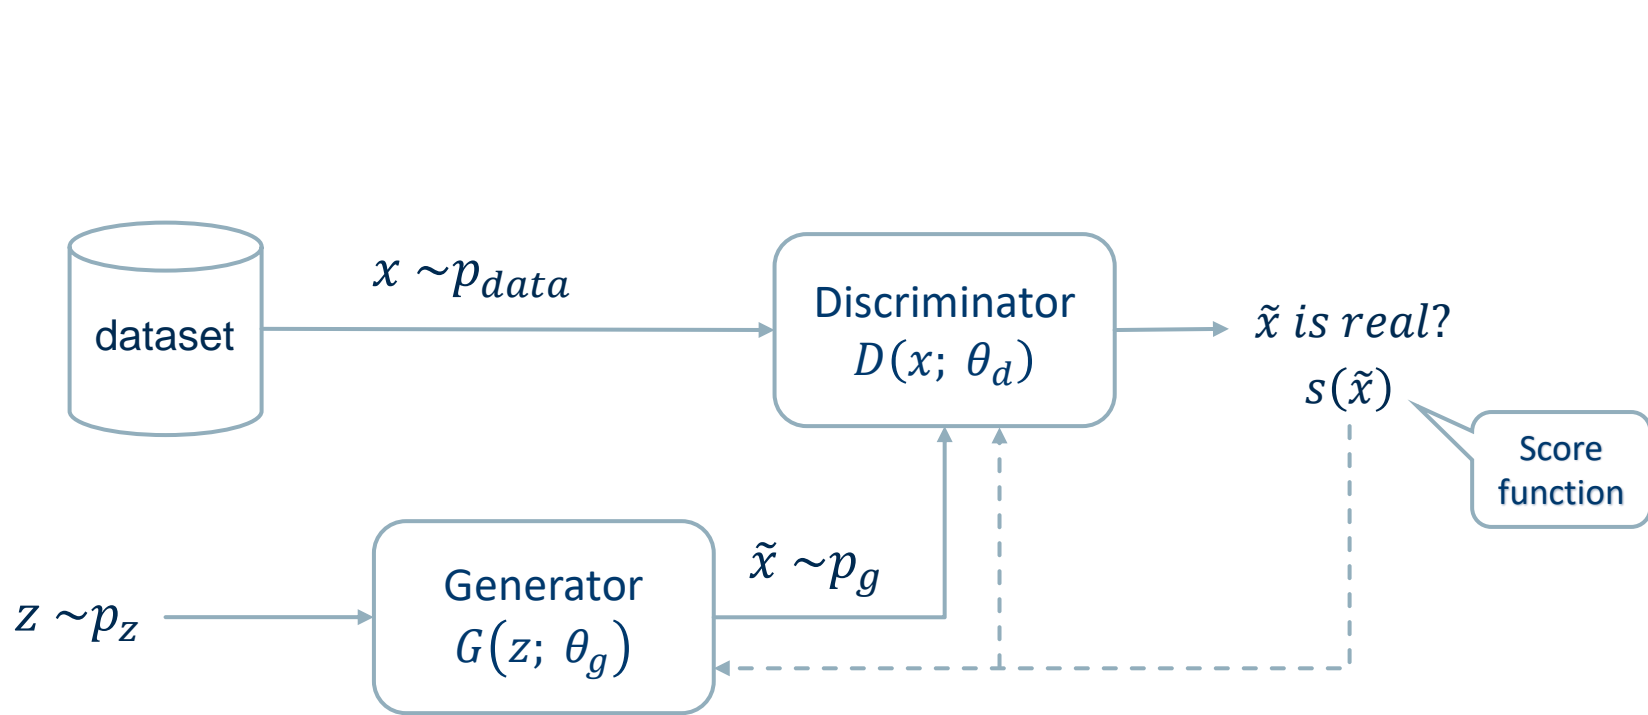
- So far: RNNs and CNNs
- Recurrent Neural Networks (RNNs) and Convolutional Neural Networks (CNNs) represent probably the biggest contribution of deep learning to the field of neural networks.
- However, deep learning is responsible for other innovations, such as for example Generative Adversarial Networks (GANs).

- GANs include two neural networks competing with each other: the generator and the discriminator.
- A **generator  $G$**  is a transformation that transforms the input noise  $z$  into a tensor – usually an image –  $x$  ( $x=G(z)$ ). The generated image  $x$  is then fed into the discriminator network  $D$ .
- The **discriminator network  $D$**  compares the real images in the training set and the image generated by the generator network and produces an output  $D(x)$ , which is the probability that image  $x$  is real.



- Both generator and discriminator are trained using the backpropagation algorithm to produce  $D(x)=1$  for the generated images  $x$ .
- Both networks are trained in alternating steps, competing with each other to improve themselves.
- The GAN model eventually converges and produces images that look real.
- Given a training set, this technique learns to generate new data under the same statistics as the training set.

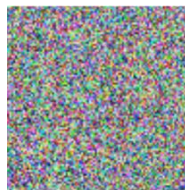
## GAN: Architecture



# GANs

- For example, a GAN trained on photographs can generate new photographs that look at least superficially authentic to human observers, having many realistic characteristics.
- GANs have been successfully applied to image tensors to create anime, human figures, and even van Gogh-like masterpieces.

Noise  $\sim N(0,1)$



Generative  
Model

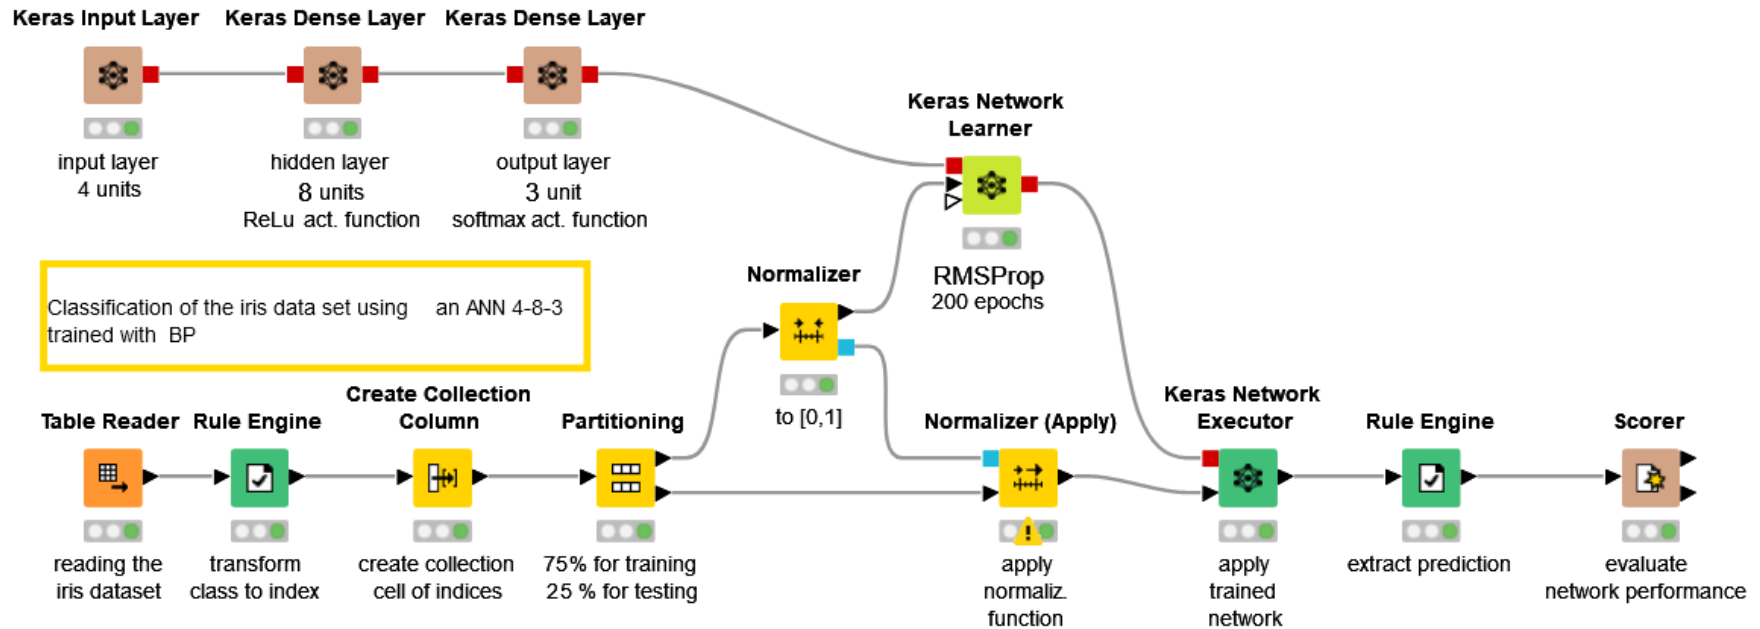


Image from: Pankaj Kishore, Towards data Science

<https://towardsdatascience.com/art-of-generative-adversarial-networks-gan-62e96a21bc35>

# Practical Example

- A multilayer perceptron with layers (4–8–3) is trained to classify the iris data set using the backpropagation algorithm, as set in the Keras Network Learner node



# Thank you

For any questions please contact: [education@knime.com](mailto:education@knime.com)