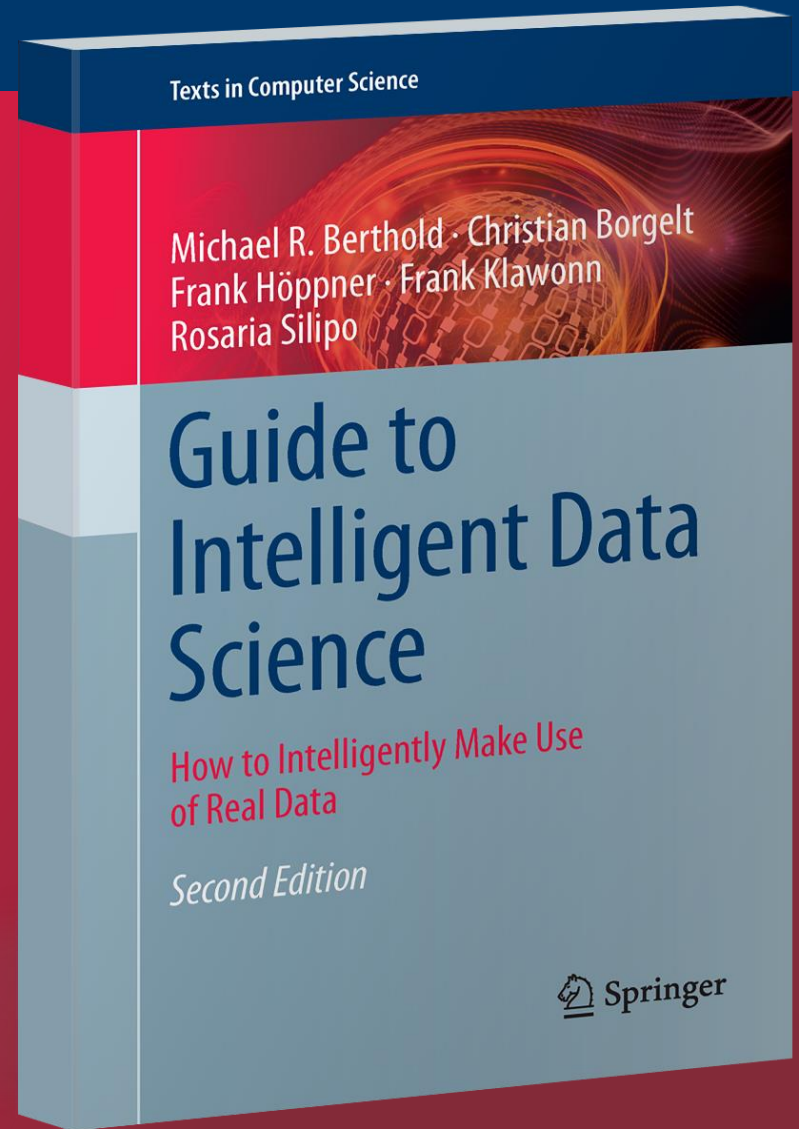


Ensemble Learning



*“If I do not believe the news in today's paper, I buy 100 copies of the paper.
Then I believe.”
-Ludwig Wittgenstein*

How can we learn from multiple models together?

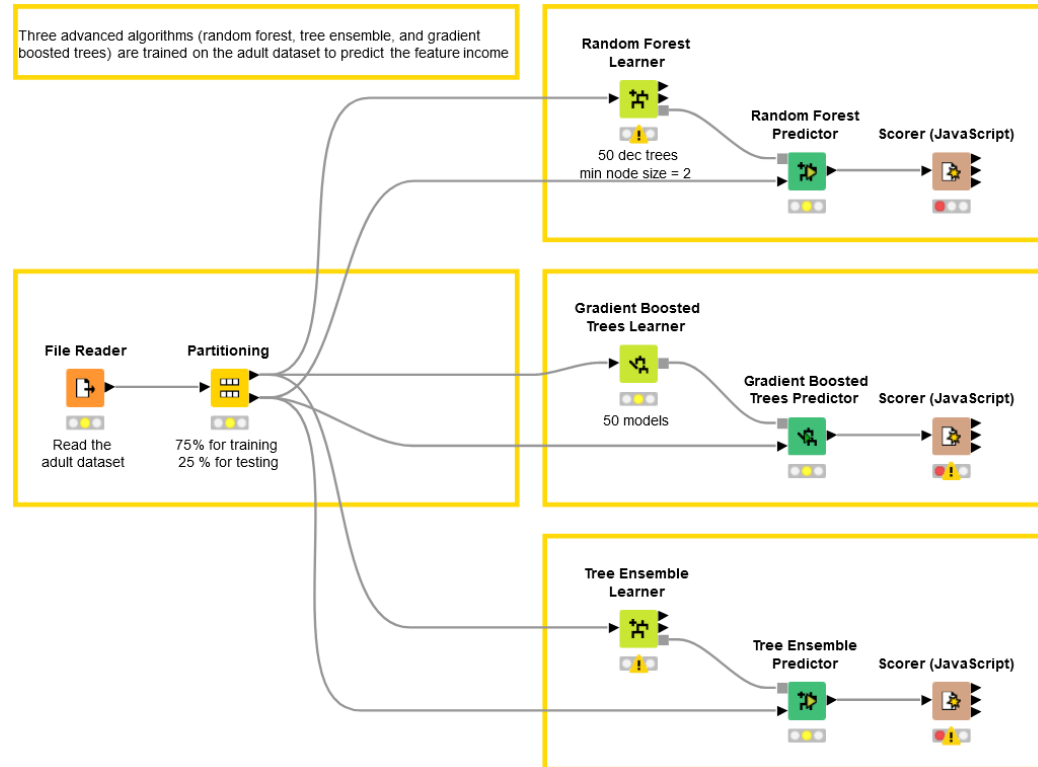
**This lesson refers to chapter 9 of the GIDS book*

- Ensemble Learning
 - Wisdom of the Crowd
 - Bagging & Boosting
 - Stacking and Cascade Generalization
 - Cascading and Delegating
 - Tree Ensembles and Random Forest
 - AdaBoost
 - Gradient Boosted Trees
 - Practical Examples

– Datasets used : adult dataset


– „Random Forest, Gradient Boosted Trees, and Tree Ensemble“ <https://kni.me/w/Ueq3QR9hty8Osh2E>

- Random forest
- Gradient boosting
- Tree ensemble



- General idea: take advantage of the “wisdom of the crowd”
- Training of many weak classifiers (or regression models)
- Combining them to construct a classifier (regression model) more accurate than any of the individual ones
- Leads to a more accurate and robust model
- Interpretation of an ensemble learning model is difficult
 - Since it consists of many models!

Wisdom of the Crowd



Crowd wiser than **any individual**


- When?
- For which questions?

- The **collective knowledge** of a *diverse* and *independent* body of people typically **exceeds** the knowledge of **any single individual** and can be harnessed by voting.
- <http://www.csc.kth.se/utbildning/kth/kurser/DD2431/ml11/schedule/07-ensemble.pdf>

Wisdom of the Crowd: Scenario


- Ask each person in the crowd:
- Will Mr. X win the general election in country Y?

- **The Crowd's prediction:**
- MAJORITY answer.
- This crowd predicts **No**. (Mr. X will not win the election.)




Wisdom of the Crowd: Scenario

- Has the crowd made a good prediction?
- Composition of crowd:
 - 30% EXPERTS.
 - 70% NON-EXPERTS.
- and their level of expertise:
 - $P(\text{correct predict}|\text{expert}) = p_e$
 - $P(\text{correct predict}|\text{non-expert}) = p_{ne}$



Wisdom of the Crowd: Scenario


- Let $p_e = 0.8$ and $p_{ne} = 0.5$
- For a random person from the crowd
- $P(\text{correct predict}|\text{individual}) = 0.3 p_e + 0.7 p_{ne} = 0.59$



Wisdom of the Crowd: Scenario

- Let $p_e = 0.8$ and $p_{ne} = 0.5$
- For a random person from the crowd
- $P(\text{correct predict}|\text{individual}) = p_i = 0.59$
- **If crowd contains 50 *independent* people:**
- $P(\text{correct predict}|\text{crowd})$

$$= \sum_{k=26}^{50} \binom{50}{k} p_i^k \cdot (1 - p_i)^{50-k} = 0.8745$$



This crowd has made a prediction with probability .875 of being correct which is $> p_i$.

It is wiser than each of the experts!


- Ensemble of predictors often outperform individual predictors
- Consider a majority voting of 5 independent classifiers in a binary classification problem.
- Each predictor, the error probability is 0.3
- Probability of three or more predictors yielding a wrong result (i.e., the majority misclassifies) is very low:

$$\sum_{i=3}^5 \binom{5}{i} 0.3^i \cdot 0.7^{5-i} = 0.08748$$


- Substantial reduction in error rate!
- In reality, classifiers are rarely independent of each other

- **Why didn't I just asked a bunch of experts??**
- Large enough crowd
 - => high probability that a sufficient number of experts will be in crowd (for any question).
- Random selection
 - => don't make a biased choice in experts.
- For some questions it may be hard to identify a diverse set of experts

- Given a **random question** expect each **person** to have a **different level of expertise**.
- Will it rain tomorrow?
 - redness proportional to expertise



- Given a **random question** expect each **person** to have a **different level of expertise**.
- Will the world go down in December?
 - redness proportional to expertise



According to *James Surowiecki* there are four elements required to form a wise crowd:

- **Diversity of opinion.** People in crowd should have a range of experiences, education and opinions. (Encourages independent predictions)
- **Independence.** Prediction by person in crowd is not influenced by other people in the crowd.
- **Decentralization.** People have specializations and local knowledge.
- **Aggregation.** There is a mechanism for aggregating all predictions into one single prediction.

In the analysis of the crowd it is implicitly assumed:

- each person is not concerned with the opinions of others, no-one is copying anyone else in the crowd.

In the analysis of the crowd we implicitly assumed:

- The non-experts will predict a **completely random wrong answer** - these will cancel each other out (to some degree).

However, there may be a systematic and consistent bias in the non-experts' predictions.


... Back to Ensemble Models

We will exploit **Wisdom of crowd** ideas for specific tasks by:

- combining (classifier) predictions
- aim to combine independent and diverse predictors (classifiers).

We can also use labeled training data

- to identify the expert classifiers in the pool;
- to identify complementary classifiers;
- to indicate how to best combine them.



Why Do Ensemble Methods Work?

Statistical reason:

- Able to average many good models
- Reduces the influence of bad models

Computational reason:

- Able to explore the model space efficiently

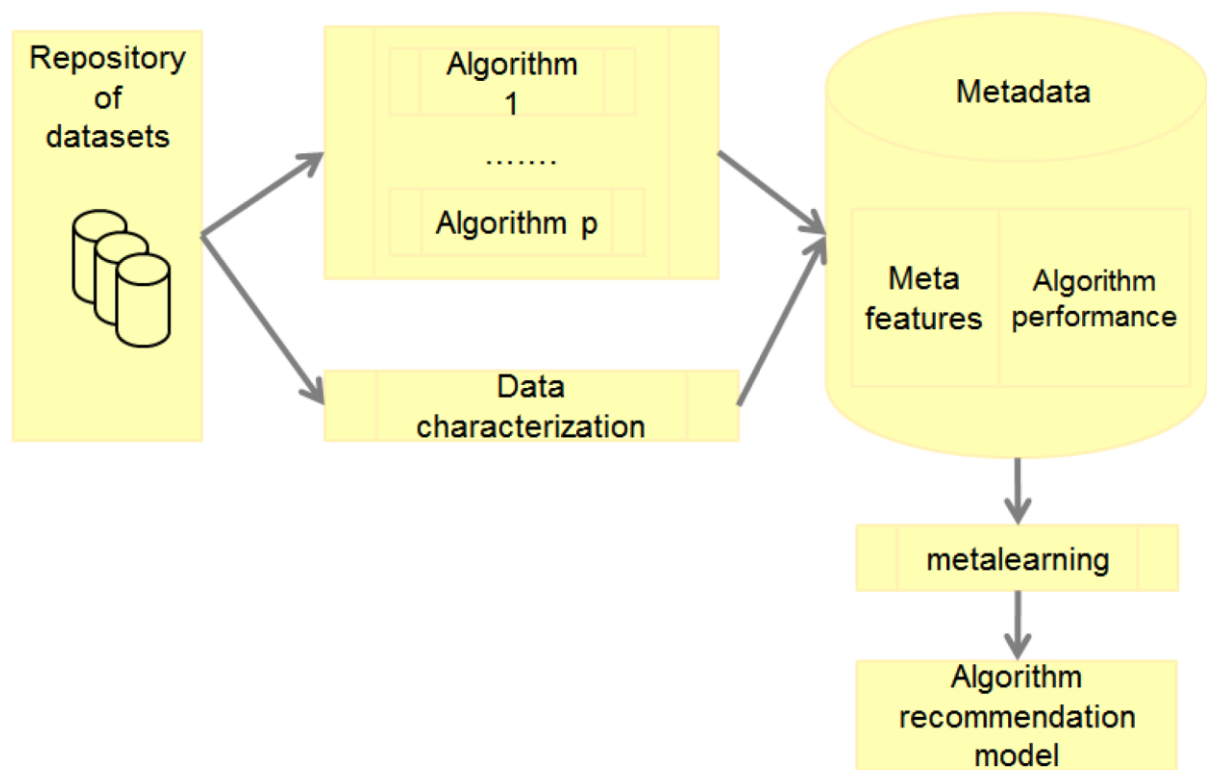
Representational reason:

- Reduce the bias of a learning algorithm by extending its model space

Remember?

- Bias = model error + algorithmic error
 - **Model error**: the error we get by selecting a model
 - **Algorithmic error**: by selecting the algorithm itself and the parameters of the algorithm
- Base-Learning: Fixed Bias / User parameterized
- Meta-Learning: Dynamic bias selection using meta-knowledge
- Meta-Knowledge: Knowledge achieved during the learning process

Ensemble Learning for Algorithm Recommendation



Combining base-learners: Categories

Philosophy	Technique
Bagging, Boosting	Variation in data
Stacking	Variation among learners (multi-expert)
Cascading, Delegating	Variation among learners (multi-stage)
Arbitrating	Variation among learners (refereed)
Meta decision trees	Variation in data and among learners

Bagging and Boosting

- Best-known techniques
- Based on selection of multiple data sub sets
- Meta model is created by combining the base models
- Advantages:
 - Reduces overfitting
 - Most effective when the base learner is highly sensitive to data
 - Typically increases accuracy
- Disadvantages:
 - Interpretability of interpretable base learners is lost

Bagging :

- Select N independent samples of the Training Data
- Learn one model on each of the samples $\Rightarrow h_1, \dots, h_N$
- Classification: Use the class most predicted by all classifiers
- Regression: Use the mean of all predictions

Boosting:

- Tries to learn a weighting for the models
- Later base learners focus more on the examples that previous base learners misclassified
- There is no single “best” boosting method

One boosting method after Schapire

Training :

- Create c_1 : base learner on a sample t_1 of the data
- Create t_2 : sample which is 50% misclassified by c_1
- Create c_2 : base learner on the sample t_2
- Create t_3 : subset of the data where c_1 predicts differently than c_2
- Create c_3 : base learner on the sample t_3

Classification :

- Classify with c_1 and c_2
- If unequal, use c_3 as final classification

Stacking and Cascade Generalization

Stacking

- In Bagging and Boosting: we used always the same base learner
- Stacking exploits differences among base learners
- Two levels of learning
 1. Base learners are trained, each on the whole data set
 2. Meta learners are created on meta data (e.g. predicted class) obtained in level 1
- Two levels of classifying
 1. Base learner are used on data point
 2. Meta learners are applied on base learner predictions


Stacking: Base learners are used in parallel

Cascade Generalization

- Base learners are used in a sequence with "partial" meta-learners
- Knowledge from previous classifiers can be used in later ones
- After each base learner has been trained, the data set is adjusted using the new information

For classification :

- Only the last model is used, which incorporates the knowledge from previous models (all base methods are used)




Cascading and Delegating

- Until now: all base classifiers are used for classification
- Here: Multistage classifiers, not all are required for classification
- Main advantage: faster classification

Cascading

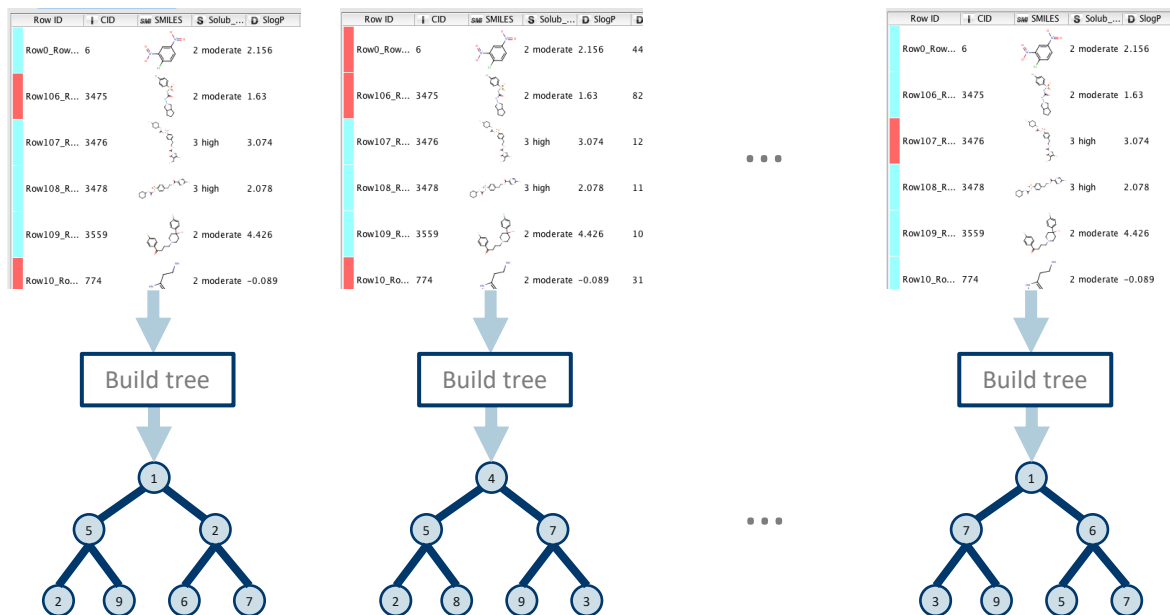
- Multilearner version of boosting
- Uses learned confidence of previous models
- Train base learner h_i using knowledge from previous base learner...
- ...on data, which was most probably misclassified by previous learners
- **Classification:** go through all base models, stop and use as classification if the model has confidence greater than epsilon



- Cascading: all instances are used in each step
- **Delegating**: only instances below confidence threshold are processed in the next step
- Idea:
 - Use everything and test for which data points you are good enough
 - Pass the remaining work to someone else.
 - If there is no someone else, ... guess
- Advantages:
 - Still understandable (no model combination)
 - Improved efficiency, due to the decreasing number of examples.

Tree Ensembles and Random Forest

- Bagging \equiv Bootstrap AGGregation
- For each tree / model a training set is generated by sampling uniformly with replacement from the standard training set



Example for Bagging

Full training set

RowID	x_1	x_2	y
Row_1	2	6	Class 1
Row_2	4	1	Class 2
Row_3	9	3	Class 2
Row_4	2	7	Class 1
Row_5	8	1	Class 2
Row_6	2	6	Class 1
Row_7	5	2	Class 2

Sampled training set

RowID	x_1	x_2	y
Row_3	9	3	Class 2
Row_6	2	6	Class 1
Row_1	2	6	Class 1
Row_3	9	3	Class 2
Row_5	8	1	Class 2
Row_6	2	6	Class 1
Row_1	2	6	Class 1


Why sampling with replacement?

- So that a sample approximates the distribution of the population
 - Frequent values are represented more
 - Less frequent values are represented less

Ultimately leads to the model with smaller variance and smaller bias

An Extra Benefit of Bagging: Out of Bag Estimation

- Able to evaluate the model using the training data
- Apply trees to samples that haven't been used for training



Out-of-bag error estimates - 0:105 - Random Forest Learner


File Hilite Navigation View

Table "default" - Rows: 2666 Spec - Columns: 26 Properties Flow Variables

Row ID	S	D	P (Churn=0)	D	P (Churn=1)	S	Churn (Out-of-bag)	D	Churn (Out-of-bag)	I	model count
Row1_Row0	S	0.943	0.057	0	0	0.943		35			
Row2_Row1	JH	1	0	0	0	1		33			
Row3_Row2	J	1	0	0	0	1		37			
Row4_Row3	JH	0.528	0.472	0	0	0.528		36			
Row5_Row4	K	0.976	0.024	0	0	0.976		41			
Row6_Row5	L	0.848	0.152	0	0	0.848		33			
Row7_Row6	4A	0.833	0.167	0	0	0.833		36			
Row9_Row8	A	0.667	0.333	0	0	0.667		30			
Row11_Row...	V	0.138	0.862	1	0	0.862		29			
Row13_Ro...	A	0.974	0.026	0	0	0.974		39			
Row14_Ro...	4T	0.917	0.083	0	0	0.917		36			
Row15_Ro...	A	0.387	0.613	1	0	0.613		31			
Row18_Ro...	T	0.974	0.026	0	0	0.974		39			
Row19_Ro...	A	1	0	0	0	1		38			
Row21_Ro...	L	0.971	0.029	0	0	0.971		34			
Row22_Ro...	O	0.03	0.97	1	0	0.97		33			
Row23_Ro...	Z	0.854	0.146	0	0	0.854		41			
Row25_Ro...	A	0.973	0.027	0	0	0.973		37			
Row26_Ro...	IE	0.886	0.114	0	0	0.886		35			
Row27_Ro...	Y	0.912	0.088	0	0	0.912		34			
Row28_Ro...	4T	0.976	0.024	0	0	0.976		42			
Row29_Ro...	4O	1	0	0	0	1		42			
Row30_Ro...	II	1	0	0	0	1		40			
Row32_Ro...	IH	0.914	0.086	0	0	0.914		35			
Row33_Ro...	A	0.875	0.125	0	0	0.875		32			
Row34_Ro...	A	0.875	0.125	0	0	0.875		32			


Random Forest

- Bag of decision trees, with an extra element of randomization
- Each node in the decision tree only “sees” a subset of the input features → ***Random Subspace Selection***
- typically \sqrt{n} to pick from
- Random forests tend to be very robust w.r.t. overfitting



AdaBoost

- Freund & Schapire (1995)
- AB is a linear classification algorithm
- AB has good generalization properties (Avoids overfitting as long as the training data is not too noisy)
- AB is a feature selector



$$\text{Strong classifier} = (\alpha_1 h_1) + (\alpha_t h_t) + \dots + (\alpha_T h_T)$$

- AdaBoost classifier: $\sum_{t=1}^T (\alpha_t h_t(x))$
- Where a weak classifier $h_t(x)$ is weighted by α_t for steps up to T
- Misclassified data points are weighted more in subsequent steps
- Classification result: $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$

- Strong Classifier: $\sum_{t=1}^T (\alpha_t h_t(\mathbf{x}))$
- Where $h_t(\mathbf{x})$ is a base classifier weighted by α_t
- Classification result: $H(\mathbf{x}) = \text{sgn}(\sum_{t=1}^T (\alpha_t h_t(\mathbf{x})))$

- Initially, all data points are given the same weight $w_{i,1} = 1/n$.
- At step t , the classifier weight α_t is calculated as

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - e_t}{1 + e_t} \right)$$

- Where e_t is given by

$$e_t = \frac{\sum_{i=1}^n w_{i,t} y_i h_t(\mathbf{x}_i)}{\sum_{i=1}^n w_{i,t}}$$

- Weights w 's are updated as

$$w_{i,t+1} = \frac{w_{i,t} \exp(-\alpha_t y_i h_t(\mathbf{x}_i))}{\sum_{j=1}^n w_{j,t} \exp(-\alpha_t y_j h_t(\mathbf{x}_j))}$$

Algorithm Ada Boost

Initialize weight of x_i with $D_0(i) = \frac{1}{m}$

for $t = 1, \dots, T$: **do**

1. $h_t = \text{"Base-Learner"}$

Calculate error $\epsilon_t = \sum_{i=1}^m D_t(i) \delta_{(y_i, h_t(x_i))}$

3. Calculate weight of learner $\alpha_t = \log \frac{1-\epsilon_t}{\epsilon_t}$

4. Update the weights $D_t(i)$ of all x_i

end for

Resulting classifier

$$H(x) = \text{sgn} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

Algorithm Base learner

Train a set H of many many base learners h on the data

Return the one h with the lowest weighted classification error

$$h = \arg \min_{h_j \in H} \epsilon = \arg \min_{h_j \in H} \left(\sum_{i=1}^m D(i) \delta_{(y_i, h_t(x_i))} \right)$$

Make sure $\epsilon < 0.5$

Algorithm Ada Boost

Initialize weight of x_i with $D_0(i) = \frac{1}{m}$

for $t = 1, \dots, T$: **do**

1. $h_t = \text{"Base-Learner"}$

Calculate error $\epsilon_t = \sum_{i=1}^m D_t(i) \delta_{(y_i, h_t(x_i))}$

3. Calculate weight of learner $\alpha_t = \log \frac{1-\epsilon_t}{\epsilon_t}$

4. Update the weights $D_t(i)$ of all x_i

end for

Resulting classifier

$$H(x) = \text{sgn} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

Update of weights of Training Samples

- Update weights and normalize them:

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \exp(-\alpha_t y_i h_t(x_i))$$

$$Z_t(i) = \sum_{i=0}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i))$$

- Weight of correct classified examples is decreased
- Weight of incorrect classified examples is increased

- Weighting the base learners “Upper-Bound Theorem”
- Primary goal is to minimize

$$\epsilon_{tr}(H) = \frac{1}{m} |\{i: H(x_i) \neq y_i\}|$$

- Global error is bounded by

$$\epsilon_{tr}(H) \leq Z_t \quad t = 1, \dots, T$$

$$Z_t = \sum_{i=0}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i))$$

- Weighting the base learners
- That's why
- Minimizing $Z_t = \sum_{i=0}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i))$ results in a minimization of the global error
- Upper-Bound can be minimized by ...
 1. Choosing the optimal hypothesis h_t ...
 2. ... with an optimal weight α_t
- Minimizing Z_t results in $\alpha_t = \log \frac{1 - \epsilon_t}{\epsilon_t}$

Advantages

- Very simple to implement
- Feature selection on very large features spaces
- Fairly good generalization

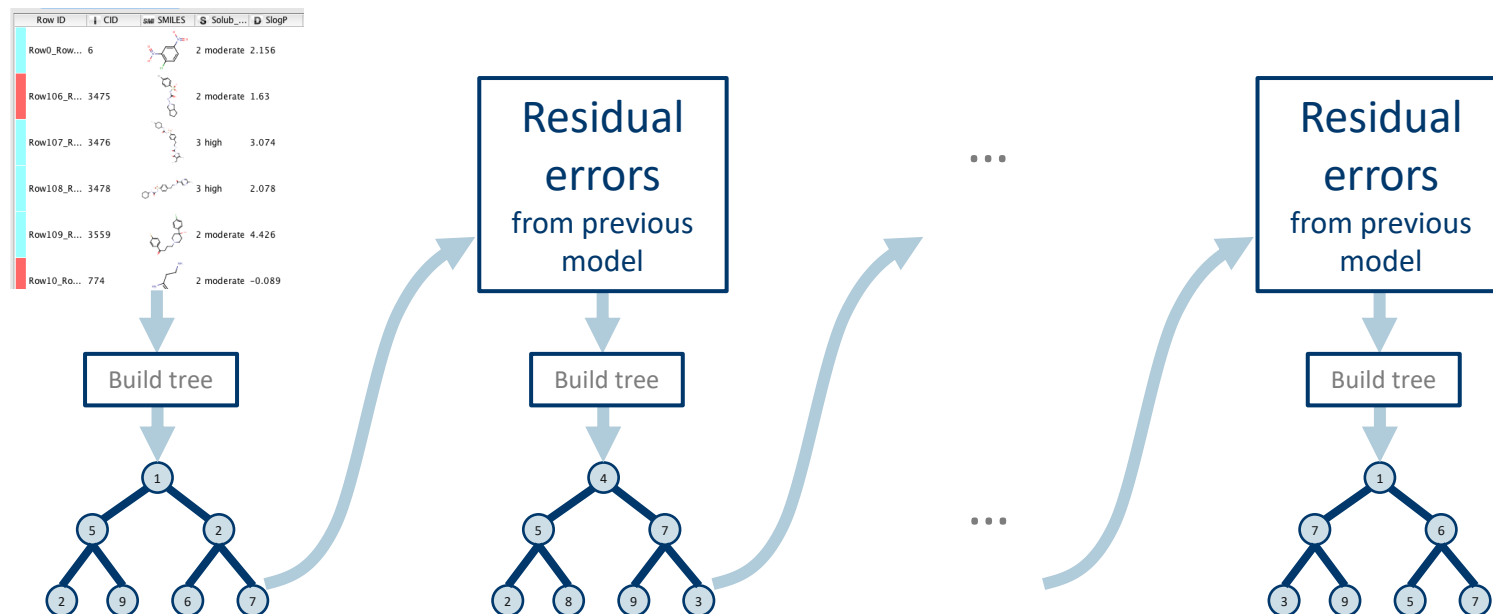
Disadvantages

- Can overfit in presence of noise
- Unclear which weak-learning algorithm fits best for a given problem

Gradient Boosted Trees

Boosting - Idea

- Starts with a single tree built from the data
- Fits a tree to residual errors from the previous model to refine the model sequentially



A shallow tree (depth 4 or less) is built at each step

- To fit residual errors from the previous step
- Resulting in a tree $h_m(x)$


The resulting tree is added to the latest model to update

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x)$$


- Where $F_{m-1}(x)$ is the model from the previous step
- The weight γ_m is chosen to minimize the loss function
 - Loss function: quantifies the difference between model predictions and data

Gradient Boosted Trees Example – Regression


Iteration 1




Iteration 5




Iteration 20




Iteration 2



Iteration 10



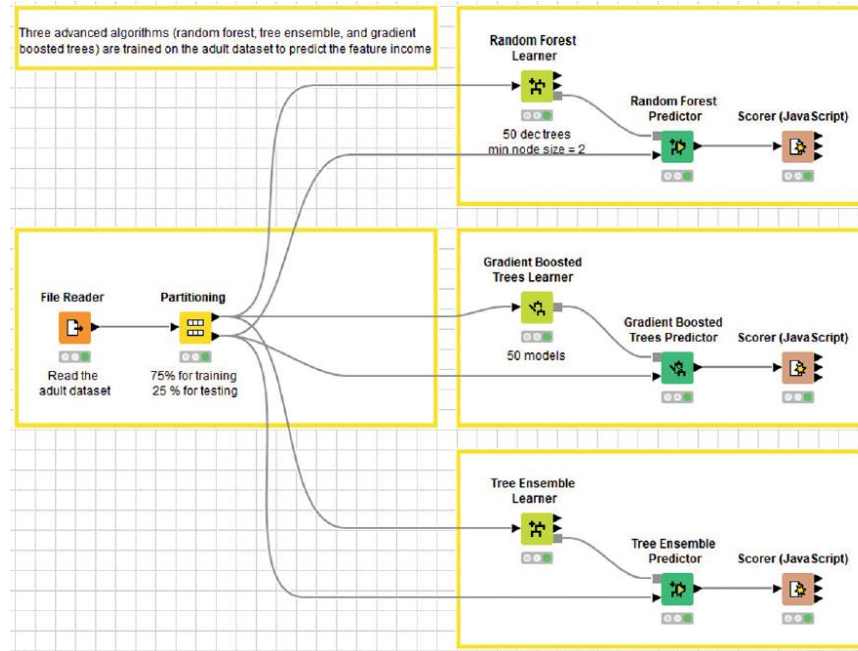
Iteration 50



- Can be used for classification and regression
- Large number of iterations – prone to overfitting
 - ~100 iterations are sufficient
- Can introduce randomness in choice of data subsets (“stochastic gradient boosting”) and choice of input features

Practical Examples with KNIME Analytics Platform

Tree Ensemble, Random Forest, and Gradient Boosted Tree



- Tree Ensemble, Random Forest, and Gradient Boosted Tree applied to the adult dataset

Thank you

For any questions please contact: education@knime.com